

```
In [1]: # Data Loading and Cleaning
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os

# Load the dataset

file_path = "D:/capstone/datasets/Affinity - State - Daily.xlsx"
df = pd.read_excel(file_path)

print(f"Raw dataset shape: {df.shape}")
print("\nFirst few rows of raw data:")
print(df.head())

# Create a proper datetime column
if 'year' in df.columns and 'month' in df.columns and 'day' in df.columns:
    print("\nCreating datetime column from year, month, day...")
    df['date'] = pd.to_datetime(df[['year', 'month', 'day']])

# Identify all spend-related columns
spend_columns = [col for col in df.columns if 'spend' in col]
print(f"\nIdentified {len(spend_columns)} spending-related columns: {spend_columns}")

# Replace '.' with NaN (missing values)
print("\nReplacing '.' with NaN in spending columns...")
df[spend_columns] = df[spend_columns].replace('.', pd.NA)

# Convert all spend columns to numeric
print("\nConverting spend columns to numeric...")
df[spend_columns] = df[spend_columns].apply(pd.to_numeric, errors='coerce')

# Check for missing values
missing_counts = df[spend_columns].isnull().sum()
print("\nMissing values count before imputation:")
print(missing_counts)

# Interpolate missing values (best-performing imputation method)
print("\nInterpolating missing values...")
df[spend_columns] = df[spend_columns].interpolate()

# Forward fill any remaining NaNs (especially at start of series)
print("\nForward filling any remaining missing values...")
df[spend_columns] = df[spend_columns].fillna(method='ffill')

# Backward fill any remaining NaNs (if any at the end of series)
df[spend_columns] = df[spend_columns].fillna(method='bfill')

# Check for any remaining missing values
remaining_missing = df[spend_columns].isnull().sum().sum()
print(f"\nRemaining missing values after imputation: {remaining_missing}")

# Drop rows where spend_all is still missing
print("\nRemoving rows where 'spend_all' is still missing...")
```

```
df_timeseries = df.dropna(subset=['spend_all'])

# Filter to include only necessary columns for our analysis
df_timeseries = df_timeseries[['date', 'spend_all']]

# Check for duplicate dates
duplicates = df_timeseries['date'].duplicated().sum()
if duplicates > 0:
    print(f"\nWarning: Found {duplicates} duplicate dates. Keeping the first occurrence")
    df_timeseries = df_timeseries.drop_duplicates(subset=['date'], keep='first')

# Sort by date to ensure chronological order
df_timeseries = df_timeseries.sort_values('date').reset_index(drop=True)

# Examine the cleaned dataset
print("\nCleaned dataset information:")
print(f"Date range: {df_timeseries['date'].min()} to {df_timeseries['date'].max()}")
print(f"Number of rows: {len(df_timeseries)}")
print(f"Number of unique dates: {df_timeseries['date'].nunique()}")

# Check for gaps in dates
print("\nChecking for gaps in the time series...")
date_range = pd.date_range(start=df_timeseries['date'].min(), end=df_timeseries['date'].max())
missing_dates = set(date_range) - set(df_timeseries['date'])
if missing_dates:
    print(f"Found {len(missing_dates)} missing dates in the time series.")
    print("Filling in missing dates...")
    # Create a complete date series
    full_date_df = pd.DataFrame({'date': date_range})
    # Merge with existing data
    df_timeseries = pd.merge(full_date_df, df_timeseries, on='date', how='left')
    # Interpolate missing spend_all values
    df_timeseries['spend_all'] = df_timeseries['spend_all'].interpolate()
else:
    print("No gaps found in the time series.")

# Summary statistics for spend_all
print("\nSummary statistics for spend_all:")
print(df_timeseries['spend_all'].describe())

# Visualize the cleaned time series
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'], df_timeseries['spend_all'])
plt.title('Consumer Spending Time Series (Cleaned)')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.grid(True)
plt.tight_layout()
plt.savefig('cleaned_timeseries.png')
plt.show()

print("\nData cleaning completed successfully. Ready for modeling.")
```

Raw dataset shape: (50694, 29)

First few rows of raw data:

	year	month	day	statefips	freq	spend_all	spend_aap	spend_acf	spend_aer	\
0	2018	12	31	1	d
1	2018	12	31	2	d
2	2018	12	31	4	d
3	2018	12	31	5	d
4	2018	12	31	6	d

	spend_apg	...	spend_sgh	spend_tws	spend_retail_w_grocery	\
0
1
2
3
4

	spend_retail_no_grocery	spend_all_incmiddle	spend_all_q1	spend_all_q2	\
0
1
2
3
4

	spend_all_q3	spend_all_q4	provisional
0	.	.	0
1	.	.	0
2	.	.	0
3	.	.	0
4	.	.	0

[5 rows x 29 columns]

Creating datetime column from year, month, day...

Identified 23 spending-related columns: ['spend_all', 'spend_aap', 'spend_acf', 'spend_aer', 'spend_apg', 'spend_durables', 'spend_nondurables', 'spend_grf', 'spend_gen', 'spend_hic', 'spend_hcs', 'spend_inperson', 'spend_inpersonmisc', 'spend_remoteservices', 'spend_sgh', 'spend_tws', 'spend_retail_w_grocery', 'spend_retail_no_grocery', 'spend_all_incmiddle', 'spend_all_q1', 'spend_all_q2', 'spend_all_q3', 'spend_all_q4']

Replacing '.' with NaN in spending columns...

Converting spend columns to numeric...

Missing values count before imputation:

spend_all	1644
spend_aap	1644
spend_acf	1644
spend_aer	1644
spend_apg	1644
spend_durables	1644
spend_nondurables	1644
spend_grf	1644
spend_gen	1644

```

spend_hic          1644
spend_hcs          1644
spend_inperson     1644
spend_inpersonmisc 1644
spend_remoteservices 1644
spend_sgh          1644
spend_tws          1644
spend_retail_w_grocery 1644
spend_retail_no_grocery 1644
spend_all_incmiddle 1644
spend_all_q1       4587
spend_all_q2       1644
spend_all_q3       1644
spend_all_q4       3606
dtype: int64

```

Interpolating missing values...

Forward filling any remaining missing values...

Remaining missing values after imputation: 0

Removing rows where 'spend_all' is still missing...

Warning: Found 49700 duplicate dates. Keeping the first occurrence...

Cleaned dataset information:

Date range: 2018-12-31 00:00:00 to 2024-06-16 00:00:00

Number of rows: 994

Number of unique dates: 994

Checking for gaps in the time series...

Found 1001 missing dates in the time series.

Filling in missing dates...

Summary statistics for spend_all:

```

count    1995.000000
mean      0.116389
std       0.132572
min      -0.321000
25%      -0.023900
50%       0.145000
75%       0.232429
max       0.352000

```

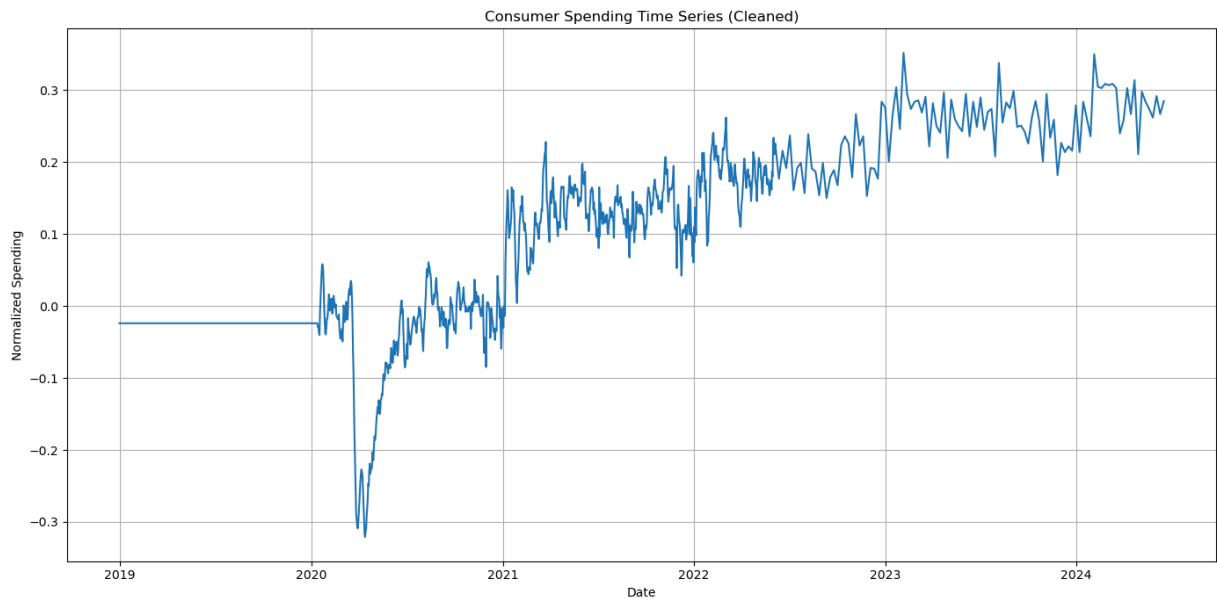
Name: spend_all, dtype: float64

C:\Users\dheer\AppData\Local\Temp\ipykernel_27932\2753603710.py:44: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
df[spend_columns] = df[spend_columns].fillna(method='ffill')
```

C:\Users\dheer\AppData\Local\Temp\ipykernel_27932\2753603710.py:47: FutureWarning: DataFrame.fillna with 'method' is deprecated and will raise in a future version. Use obj.ffill() or obj.bfill() instead.

```
df[spend_columns] = df[spend_columns].fillna(method='bfill')
```



Data cleaning completed successfully. Ready for modeling.

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import mean_absolute_error, mean_squared_error
from statsmodels.tsa.stattools import adfuller
from pmdarima import auto_arima
from prophet import Prophet
from sklearn.preprocessing import MinMaxScaler
import tensorflow as tf
from tensorflow.keras.models import Sequential, Model
from tensorflow.keras.layers import (
    LSTM, Dense, Dropout, Input, LayerNormalization,
    GlobalAveragePooling1D, MultiHeadAttention, Add
)
from tensorflow.keras.callbacks import EarlyStopping

# Set random seeds for reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# Initialize results storage for comparison
results_df = pd.DataFrame(columns=['Model', 'Data Type', 'MAE', 'MSE', 'RMSE'])

# Create a function to add results to our comparison dataframe
def add_result(model_name, data_type, mae, mse, rmse):
    global results_df
    results_df = pd.concat([results_df, pd.DataFrame({
        'Model': [model_name],
        'Data Type': [data_type],
        'MAE': [mae],
        'MSE': [mse],
        'RMSE': [rmse]
    })], ignore_index=True)
```

```

# Split data - use last 30 days for testing/evaluation
test_size = 30
train_data = df_timeseries.iloc[:-test_size].copy()
test_data = df_timeseries.iloc[-test_size:].copy()

# Also create recovery phase dataset (from 2021-01-01 onwards)
recovery_data = df_timeseries[df_timeseries['date'] >= '2021-01-01'].copy()
recovery_train = recovery_data.iloc[:-test_size].copy()
recovery_test = recovery_data.iloc[-test_size:].copy()

print(f"Full dataset: {len(df_timeseries)} days from {df_timeseries['date'].min()}")
print(f"Training data: {len(train_data)} days")
print(f"Test data: {len(test_data)} days")
print(f"Recovery phase: {len(recovery_data)} days from {recovery_data['date'].min()}")

# Visualize the data
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'], df_timeseries['spend_all'], label='Full Dataset')
plt.plot(recovery_data['date'], recovery_data['spend_all'], label='Recovery Phase (')
plt.axvline(x=pd.to_datetime('2021-01-01'), color='red', linestyle='--', label='Rec
plt.axvline(x=test_data['date'].iloc[0], color='green', linestyle='--', label='Test
plt.title('Consumer Spending Data Overview')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('data_overview.png')
plt.show()

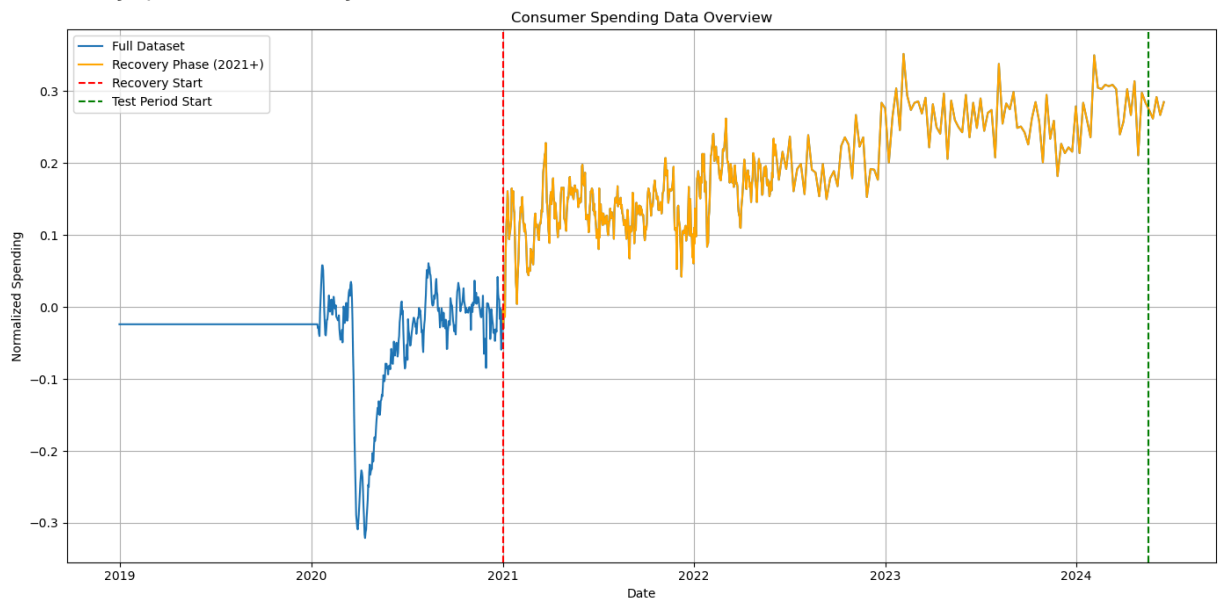
```

Full dataset: 1995 days from 2018-12-31 00:00:00 to 2024-06-16 00:00:00

Training data: 1965 days

Test data: 30 days

Recovery phase: 1263 days from 2021-01-01 00:00:00 to 2024-06-16 00:00:00



MODEL 1: ARIMA

```
In [4]: print("\n" + "="*80)
print("MODEL 1: ARIMA - FULL DATASET")
print("="*80)

# Check stationarity with ADF test
print("\nRunning ADF Test for stationarity...")
adf_result = adfuller(df_timeseries['spend_all'])
print(f"ADF Statistic: {adf_result[0]}")
print(f"p-value: {adf_result[1]}")
print("Critical Values:")
for key, value in adf_result[4].items():
    print(f"\t{key}: {value}")

# Interpret the result
if adf_result[1] < 0.05:
    print("The series is stationary")
else:
    print("The series is not stationary, differencing may be required")

# Create time series data
ts_full = train_data.set_index('date')['spend_all']
ts_test = test_data.set_index('date')['spend_all']

# Determine optimal ARIMA parameters using auto_arima
print("\nFinding optimal ARIMA parameters...")
arima_model = auto_arima(ts_full,
                        seasonal=False,
                        stepwise=True,
                        trace=True,
                        error_action='ignore',
                        suppress_warnings=True)

print(f"\nBest ARIMA model: {arima_model.order}")
print(arima_model.summary())

# Get forecast with confidence intervals
forecast_results = arima_model.predict(n_periods=test_size, return_conf_int=True)
arima_forecast = forecast_results[0] # Point forecasts
confidence_intervals = forecast_results[1] # Confidence intervals
forecast_index = ts_test.index # For plotting
```

```

# Evaluation metrics
mae_arma = mean_absolute_error(ts_test, arma_forecast)
mse_arma = mean_squared_error(ts_test, arma_forecast)
rmse_arma = np.sqrt(mse_arma)

print("\n--- ARIMA Model Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_arma:.6f}")
print(f"Mean Squared Error (MSE): {mse_arma:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_arma:.6f}")

# Add to results
add_result('ARIMA', 'Full Data', mae_arma, mse_arma, rmse_arma)

# Plot results
plt.figure(figsize=(14, 7))
plt.plot(ts_full.index, ts_full, label='Training Data')
plt.plot(ts_test.index, ts_test, label='Actual Spending')
plt.plot(ts_test.index, arma_forecast, label='ARIMA Forecast', color='red')
plt.fill_between(ts_test.index,
                 confidence_intervals[:, 0], # Lower bound
                 confidence_intervals[:, 1], # Upper bound
                 color='red', alpha=0.2, label='95% Confidence Interval')
plt.title('ARIMA Consumer Spending Forecast')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('arma_forecast.png')
plt.show()

# Forecast next 30 days (future prediction)
future_results = arma_model.predict(n_periods=30, return_conf_int=True)
future_forecast_arma = future_results[0] # Point forecasts
future_confidence_intervals = future_results[1] # Confidence intervals

# Generate future dates
future_dates = pd.date_range(start=df_timeseries['date'].max() + pd.Timedelta(days=
# Plot future forecast
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'][-90:], df_timeseries['spend_all'][-90:], label='Hist
plt.plot(future_dates, future_forecast_arma, label='ARIMA 30-Day Forecast', color=
plt.fill_between(future_dates,
                 future_confidence_intervals[:, 0], # Lower bound
                 future_confidence_intervals[:, 1], # Upper bound
                 color='red', alpha=0.2, label='95% Confidence Interval')
plt.title('ARIMA 30-Day Consumer Spending Forecast')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('arma_future_forecast.png')
plt.show()

```


=====

MODEL 1: ARIMA - FULL DATASET

=====

Running ADF Test for stationarity...

ADF Statistic: -1.3702487050072258

p-value: 0.5964532649851969

Critical Values:

1%: -3.4336771595431106

5%: -2.863009746829746

10%: -2.5675524325901415

The series is not stationary, differencing may be required

Finding optimal ARIMA parameters...

Performing stepwise search to minimize aic

```
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-11869.298, Time=0.30 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-11809.837, Time=0.20 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-11865.823, Time=0.20 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-11857.550, Time=0.78 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=-11811.515, Time=0.09 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-11874.024, Time=0.21 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-11869.555, Time=0.48 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-11877.603, Time=0.77 sec
ARIMA(2,1,1)(0,0,0)[0] intercept : AIC=-11861.822, Time=1.46 sec
ARIMA(2,1,0)(0,0,0)[0] intercept : AIC=-11876.356, Time=0.44 sec
ARIMA(1,1,1)(0,0,0)[0]          : AIC=-11872.614, Time=0.13 sec
```

Best model: ARIMA(1,1,1)(0,0,0)[0] intercept

Total fit time: 5.054 seconds

Best ARIMA model: (1, 1, 1)

SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:          1965
Model:                SARIMAX(1, 1, 1)  Log Likelihood          5942.802
Date:                 Sat, 29 Mar 2025  AIC                  -11877.603
Time:                 00:21:47         BIC                  -11855.272
Sample:               12-31-2018       HQIC                  -11869.396
                   - 05-17-2024
```

Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept    6.723e-05      0.000        0.392      0.695      -0.000      0.000
ar.L1         0.5257        0.065        8.042      0.000        0.398      0.654
ma.L1        -0.3615        0.068       -5.291      0.000       -0.495     -0.228
sigma2         0.0001     2.28e-06       60.337      0.000        0.000      0.000
=====
```

```
=====
Ljung-Box (L1) (Q):          0.19  Jarque-Bera (JB):          6195.21
Prob(Q):                    0.66  Prob(JB):              0.00
Heteroskedasticity (H):      0.51  Skew:                  0.57
Prob(H) (two-sided):         0.00  Kurtosis:              11.63
=====
```

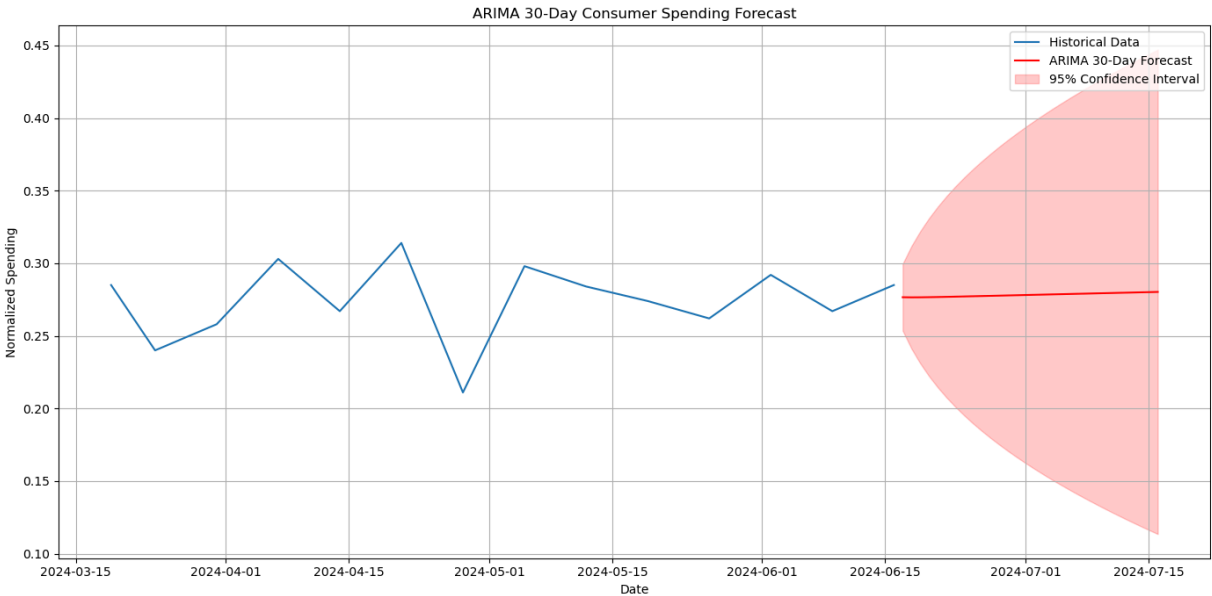
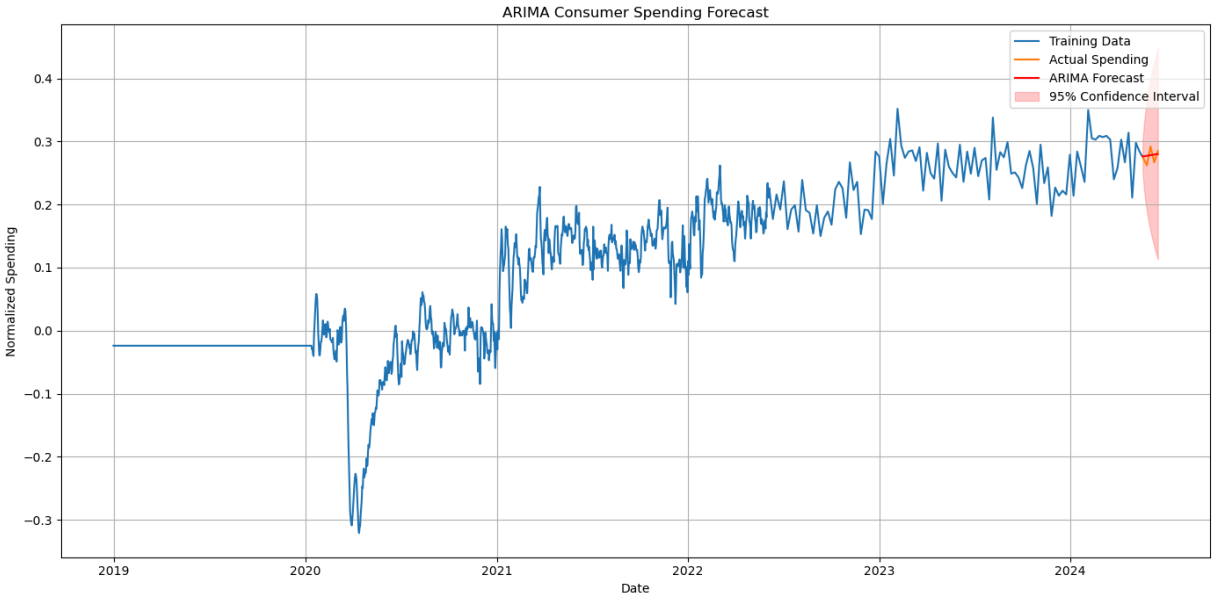
Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-ste

p).

```
--- ARIMA Model Evaluation ---  
Mean Absolute Error (MAE): 0.006686  
Mean Squared Error (MSE): 0.000062  
Root Mean Squared Error (RMSE): 0.007886
```

```
C:\Users\dheer\AppData\Local\Temp\ipykernel_27932\864976347.py:30: FutureWarning: Th  
e behavior of DataFrame concatenation with empty or all-NA entries is deprecated. In  
a future version, this will no longer exclude empty or all-NA columns when determini  
ng the result dtypes. To retain the old behavior, exclude the relevant entries before  
the concat operation.  
results_df = pd.concat([results_df, pd.DataFrame({
```



MODEL 1: ARIMA - RECOVERY PHASE


```
In [6]: print("\n" + "="*80)
print("MODEL 1: ARIMA - RECOVERY PHASE")
print("="*80)

# Split recovery data into train and test
recovery_train = recovery_data.iloc[:-test_size].copy()
recovery_test = recovery_data.iloc[-test_size:].copy()
print(f"Recovery training data: {len(recovery_train)} days")
print(f"Recovery test data: {len(recovery_test)} days")

# Create time series data for recovery phase
ts_recovery = recovery_train.set_index('date')['spend_all']
ts_recovery_test = recovery_test.set_index('date')['spend_all']

# Check stationarity
adf_result_recovery = adfuller(recovery_data['spend_all'])
if adf_result_recovery[1] < 0.05:
    print("Recovery phase data is stationary")
else:
    print("Recovery phase data is not stationary, differencing may be required")

# Find optimal parameters
arma_model_recovery = auto_arima(ts_recovery,
                                  seasonal=False,
                                  stepwise=True,
                                  trace=True,
                                  error_action='ignore',
                                  suppress_warnings=True)

print(f"\nBest ARIMA model for recovery data: {arma_model_recovery.order}")

# Forecast for test period with confidence intervals
forecast_recovery_results = arma_model_recovery.predict(n_periods=test_size, return
arma_forecast_recovery = forecast_recovery_results[0] # Point forecasts
recovery_confidence_intervals = forecast_recovery_results[1] # Confidence interval

# Evaluation metrics
mae_arma_recovery = mean_absolute_error(ts_recovery_test, arma_forecast_recovery)
mse_arma_recovery = mean_squared_error(ts_recovery_test, arma_forecast_recovery)
rmse_arma_recovery = np.sqrt(mse_arma_recovery)

print("\n--- ARIMA Recovery Phase Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_arma_recovery:.6f}")
print(f"Mean Squared Error (MSE): {mse_arma_recovery:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_arma_recovery:.6f}")

# Add to results
```

```
add_result('ARIMA', 'Recovery Phase', mae_arma_recovery, mse_arma_recovery, rmse_

# Plot for recovery phase
plt.figure(figsize=(14, 7))
plt.plot(ts_recovery.index, ts_recovery, label='Training Data (Recovery)')
plt.plot(ts_recovery_test.index, ts_recovery_test, label='Actual Spending')
plt.plot(ts_recovery_test.index, arma_forecast_recovery, label='ARIMA Forecast', c
plt.fill_between(ts_recovery_test.index,
                  recovery_confidence_intervals[:, 0], # Lower bound
                  recovery_confidence_intervals[:, 1], # Upper bound
                  color='red', alpha=0.2, label='95% Confidence Interval')
plt.title('ARIMA Consumer Spending Forecast (Recovery Phase)')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('arma_recovery_forecast.png')
plt.show()
```

```
=====
MODEL 1: ARIMA - RECOVERY PHASE
=====
```

Recovery training data: 1233 days

Recovery test data: 30 days

Recovery phase data is not stationary, differencing may be required

Performing stepwise search to minimize aic

```
ARIMA(2,1,2)(0,0,0)[0] intercept : AIC=-7288.769, Time=0.48 sec
ARIMA(0,1,0)(0,0,0)[0] intercept : AIC=-7258.442, Time=0.11 sec
ARIMA(1,1,0)(0,0,0)[0] intercept : AIC=-7287.416, Time=0.24 sec
ARIMA(0,1,1)(0,0,0)[0] intercept : AIC=-7282.501, Time=0.35 sec
ARIMA(0,1,0)(0,0,0)[0]          : AIC=-7260.046, Time=0.17 sec
ARIMA(1,1,2)(0,0,0)[0] intercept : AIC=-7293.887, Time=0.33 sec
ARIMA(0,1,2)(0,0,0)[0] intercept : AIC=-7292.841, Time=0.75 sec
ARIMA(1,1,1)(0,0,0)[0] intercept : AIC=-7293.092, Time=0.57 sec
ARIMA(1,1,3)(0,0,0)[0] intercept : AIC=-7292.138, Time=0.70 sec
ARIMA(0,1,3)(0,0,0)[0] intercept : AIC=-7293.887, Time=0.65 sec
ARIMA(0,1,4)(0,0,0)[0] intercept : AIC=-7294.048, Time=1.56 sec
ARIMA(1,1,4)(0,0,0)[0] intercept : AIC=-7309.851, Time=1.30 sec
ARIMA(2,1,4)(0,0,0)[0] intercept : AIC=-7353.399, Time=1.44 sec
ARIMA(2,1,3)(0,0,0)[0] intercept : AIC=-7290.689, Time=1.35 sec
ARIMA(3,1,4)(0,0,0)[0] intercept : AIC=-7345.162, Time=1.56 sec
ARIMA(2,1,5)(0,0,0)[0] intercept : AIC=-7332.011, Time=1.75 sec
ARIMA(1,1,5)(0,0,0)[0] intercept : AIC=-7325.515, Time=2.64 sec
ARIMA(3,1,3)(0,0,0)[0] intercept : AIC=-7288.536, Time=0.67 sec
ARIMA(3,1,5)(0,0,0)[0] intercept : AIC=-7385.558, Time=2.19 sec
ARIMA(4,1,5)(0,0,0)[0] intercept : AIC=-7340.414, Time=2.15 sec
ARIMA(4,1,4)(0,0,0)[0] intercept : AIC=-7359.860, Time=2.25 sec
ARIMA(3,1,5)(0,0,0)[0]          : AIC=-7421.713, Time=1.04 sec
ARIMA(2,1,5)(0,0,0)[0]          : AIC=-7387.368, Time=0.94 sec
ARIMA(3,1,4)(0,0,0)[0]          : AIC=-7363.926, Time=0.89 sec
ARIMA(4,1,5)(0,0,0)[0]          : AIC=-7355.362, Time=1.13 sec
ARIMA(2,1,4)(0,0,0)[0]          : AIC=-7374.794, Time=0.78 sec
ARIMA(4,1,4)(0,0,0)[0]          : AIC=-7361.919, Time=0.93 sec
```

Best model: ARIMA(3,1,5)(0,0,0)[0]

Total fit time: 28.951 seconds

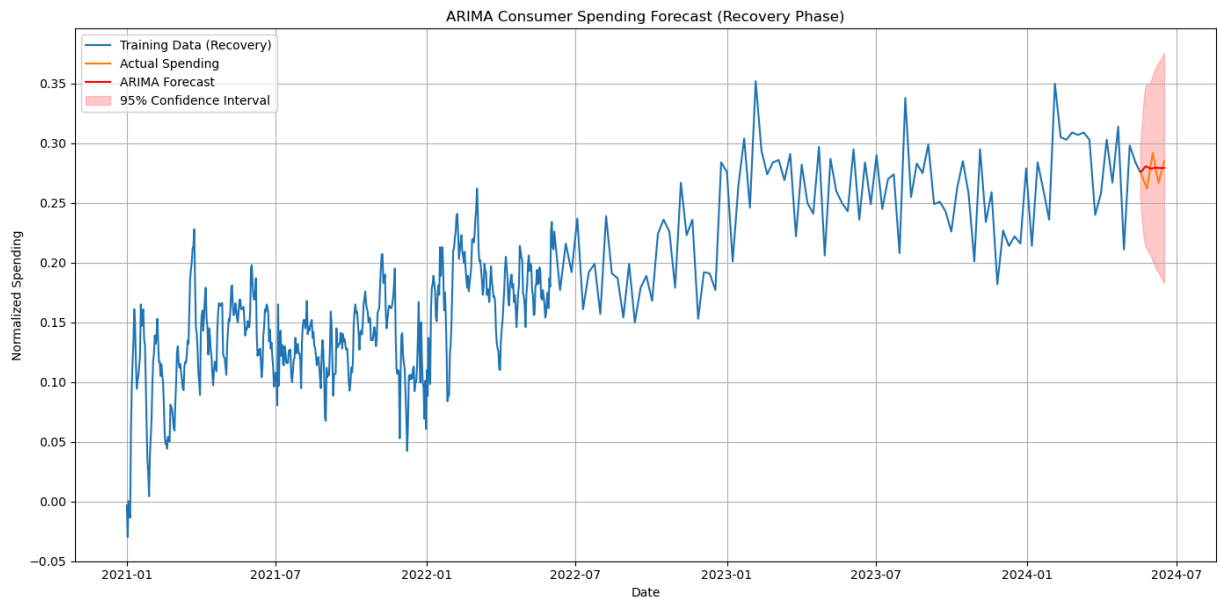
Best ARIMA model for recovery data: (3, 1, 5)

--- ARIMA Recovery Phase Evaluation ---

Mean Absolute Error (MAE): 0.007297

Mean Squared Error (MSE): 0.000078

Root Mean Squared Error (RMSE): 0.008813



MODEL 2: PROPHET

```
In [8]: print("\n" + "="*80)
print("MODEL 2: PROPHET - FULL DATASET")
print("="*80)

# Prophet requires columns named 'ds' and 'y'
prophet_train = train_data.rename(columns={'date': 'ds', 'spend_all': 'y'})
prophet_test = test_data.rename(columns={'date': 'ds', 'spend_all': 'y'})

# Initialize and train Prophet model
print("Training Prophet model with tuned parameters...")
prophet_model = Prophet(
    changepoint_prior_scale=0.05,
    seasonality_prior_scale=10.0,
    seasonality_mode='additive',
    daily_seasonality=False,
    weekly_seasonality=True,
    yearly_seasonality=True
)
prophet_model.fit(prophet_train)

# Create a future dataframe for the test period
future = prophet_model.make_future_dataframe(periods=test_size)
```

```

forecast = prophet_model.predict(future)

# Extract forecasted values for the test period
y_pred_prophet = forecast.iloc[-test_size:]['yhat'].values
y_true_prophet = prophet_test['y'].values

# Calculate evaluation metrics
mae_prophet = mean_absolute_error(y_true_prophet, y_pred_prophet)
mse_prophet = mean_squared_error(y_true_prophet, y_pred_prophet)
rmse_prophet = np.sqrt(mse_prophet)

print("\n--- Prophet Model Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_prophet:.6f}")
print(f"Mean Squared Error (MSE): {mse_prophet:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_prophet:.6f}")

# Add to results
add_result('Prophet', 'Full Data', mae_prophet, mse_prophet, rmse_prophet)

# Plot components of the forecast
fig1 = prophet_model.plot_components(forecast)
plt.tight_layout()
plt.savefig('prophet_components.png')
plt.show()

# Create a more detailed plot
plt.figure(figsize=(14, 7))
plt.plot(prophet_train['ds'], prophet_train['y'], 'ko', markersize=2, label='Train')
plt.plot(prophet_test['ds'], prophet_test['y'], 'bo', markersize=3, label='Actual')
plt.plot(forecast['ds'], forecast['yhat'], 'r-', label='Prophet Forecast')
plt.fill_between(forecast['ds'],
                 forecast['yhat_lower'],
                 forecast['yhat_upper'],
                 color='red', alpha=0.2, label='95% Confidence Interval')
plt.title('Prophet Consumer Spending Forecast with Confidence Intervals')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('prophet_detailed_forecast.png')
plt.show()

# Forecast next 30 days (future prediction)
future_forecast_df = prophet_model.make_future_dataframe( periods=test_size + 30)
future_forecast = prophet_model.predict(future_forecast_df)

# Plot future 30-day forecast
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'][-90:], df_timeseries['spend_all'][-90:], 'ko', marker=1)
plt.plot(future_forecast['ds'][-60:], future_forecast['yhat'][-60:], 'r-', label='Prophet Forecast')
plt.fill_between(future_forecast['ds'][-60:],
                 future_forecast['yhat_lower'][-60:],
                 future_forecast['yhat_upper'][-60:],
                 color='red', alpha=0.2, label='95% Confidence Interval')
future_start_idx = len(future_forecast) - 30

```

```
plt.axvline(x=future_forecast['ds'][future_start_idx-1], color='green', linestyle='dashed')
plt.title('Prophet 30-Day Consumer Spending Forecast')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('prophet_future_forecast.png')
plt.show()
```

=====

MODEL 2: PROPHET - FULL DATASET

=====

Training Prophet model with tuned parameters...

00:22:18 - cmdstanpy - INFO - Chain [1] start processing

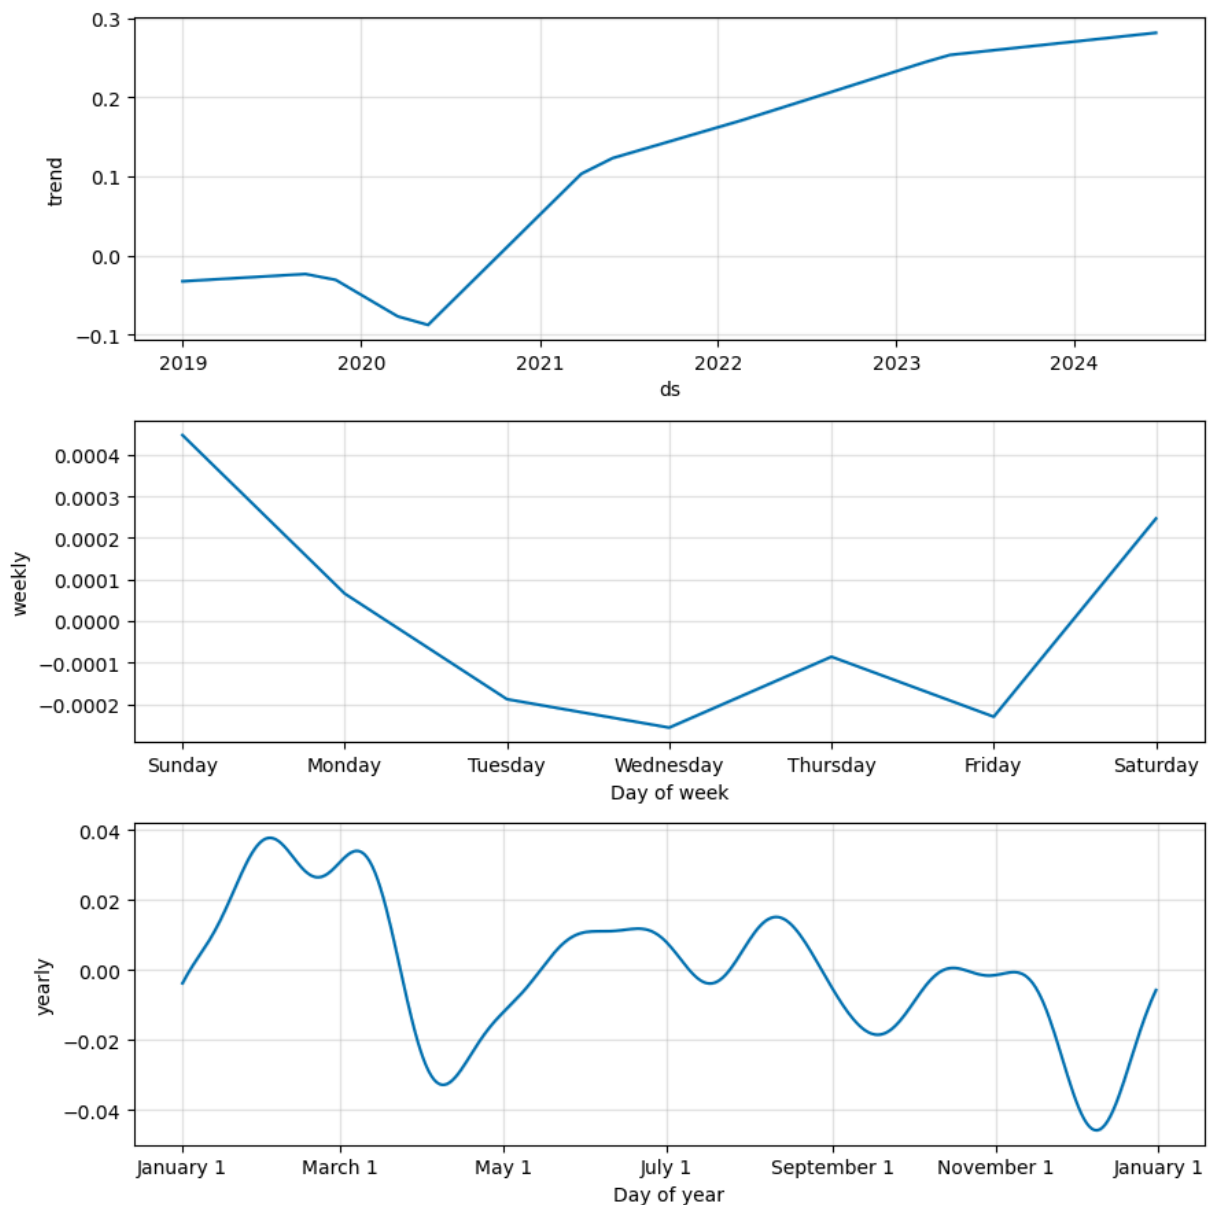
00:22:18 - cmdstanpy - INFO - Chain [1] done processing

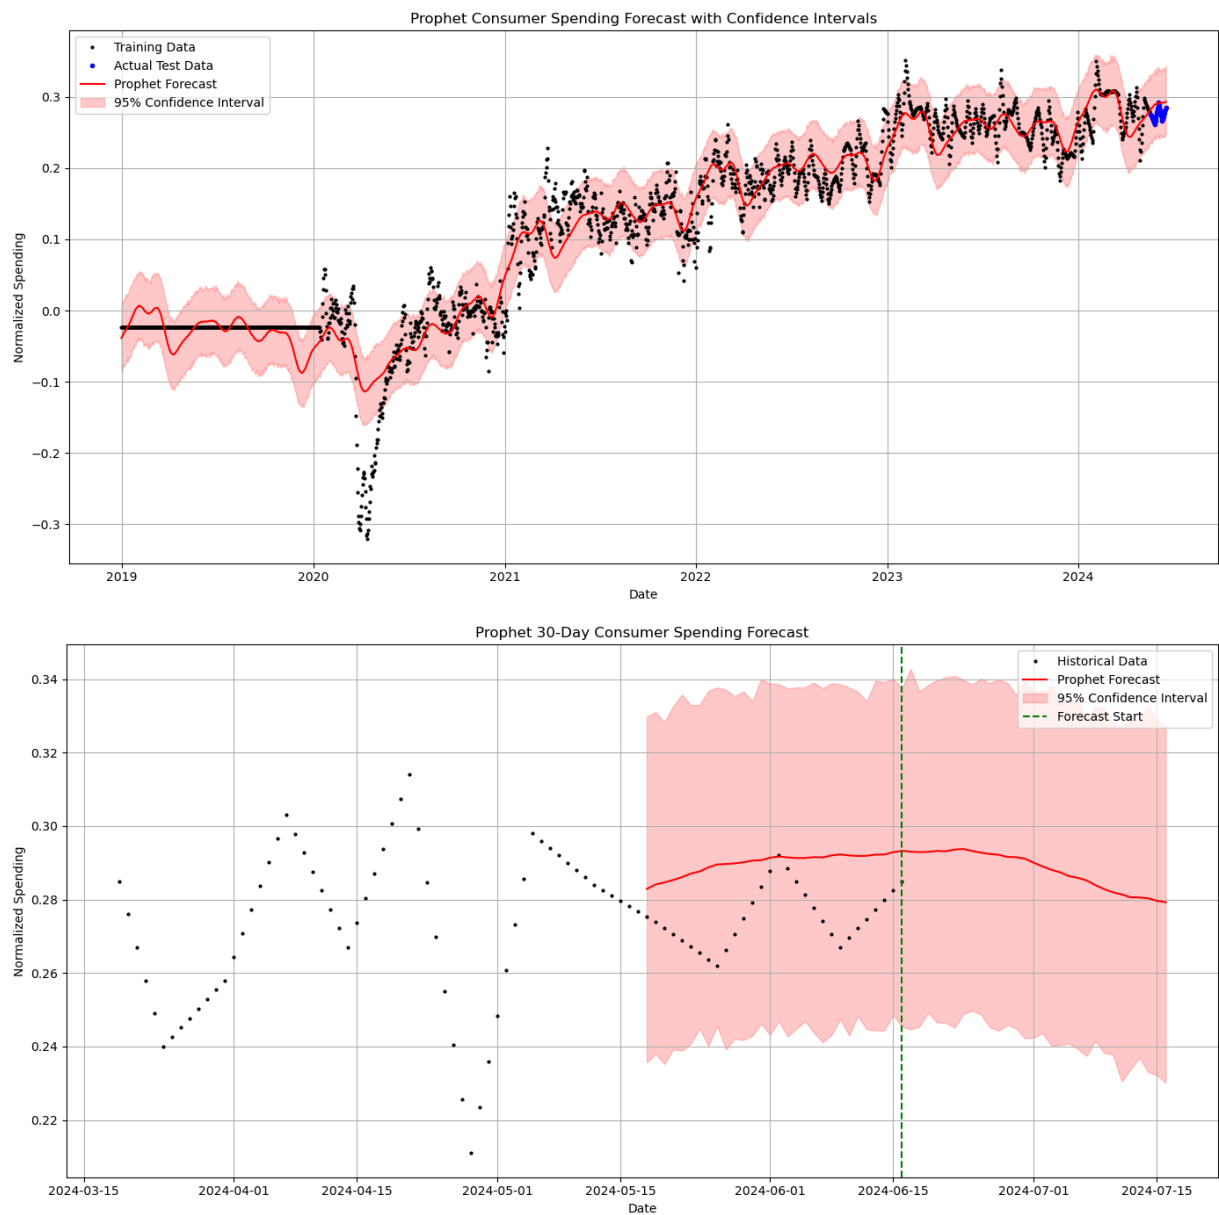
--- Prophet Model Evaluation ---

Mean Absolute Error (MAE): 0.014657

Mean Squared Error (MSE): 0.000264

Root Mean Squared Error (RMSE): 0.016253





Run Prophet on recovery phase data

```
In [10]: print("\n" + "="*80)
print("MODEL 2: PROPHET - RECOVERY PHASE")
print("="*80)
```

```

# Prepare recovery data for Prophet
prophet_recovery_train = recovery_train.rename(columns={'date': 'ds', 'spend_all': 'y'})
prophet_recovery_test = recovery_test.rename(columns={'date': 'ds', 'spend_all': 'y'})

# Train Prophet model on recovery data
prophet_model_recovery = Prophet(
    changepoint_prior_scale=0.05,
    seasonality_prior_scale=10.0,
    seasonality_mode='additive'
)
prophet_model_recovery.fit(prophet_recovery_train)

# Create future dataframe for the test period
future_recovery = prophet_model_recovery.make_future_dataframe(periods=test_size)
forecast_recovery = prophet_model_recovery.predict(future_recovery)

# Extract forecasted values for the test period
y_pred_prophet_recovery = forecast_recovery.iloc[-test_size:]['yhat'].values
y_true_prophet_recovery = prophet_recovery_test['y'].values

# Calculate evaluation metrics
mae_prophet_recovery = mean_absolute_error(y_true_prophet_recovery, y_pred_prophet_recovery)
mse_prophet_recovery = mean_squared_error(y_true_prophet_recovery, y_pred_prophet_recovery)
rmse_prophet_recovery = np.sqrt(mse_prophet_recovery)

print("\n--- Prophet Recovery Phase Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_prophet_recovery:.6f}")
print(f"Mean Squared Error (MSE): {mse_prophet_recovery:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_prophet_recovery:.6f}")

# Add to results
add_result('Prophet', 'Recovery Phase', mae_prophet_recovery, mse_prophet_recovery, rmse_prophet_recovery)

# Plot recovery results
plt.figure(figsize=(14, 7))
plt.plot(prophet_recovery_train['ds'], prophet_recovery_train['y'], 'ko', markersize=10)
plt.plot(prophet_recovery_test['ds'], prophet_recovery_test['y'], 'bo', markersize=10)
plt.plot(forecast_recovery['ds'], forecast_recovery['yhat'], 'r-', label='Prophet Forecast')
plt.fill_between(forecast_recovery['ds'],
                 forecast_recovery['yhat_lower'],
                 forecast_recovery['yhat_upper'],
                 color='red', alpha=0.2, label='95% Confidence Interval')
plt.title('Prophet Consumer Spending Forecast (Recovery Phase)')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('prophet_recovery_forecast.png')
plt.show()

```

```
00:22:20 - cmdstanpy - INFO - Chain [1] start processing
```

```
=====
MODEL 2: PROPHET - RECOVERY PHASE
=====
```

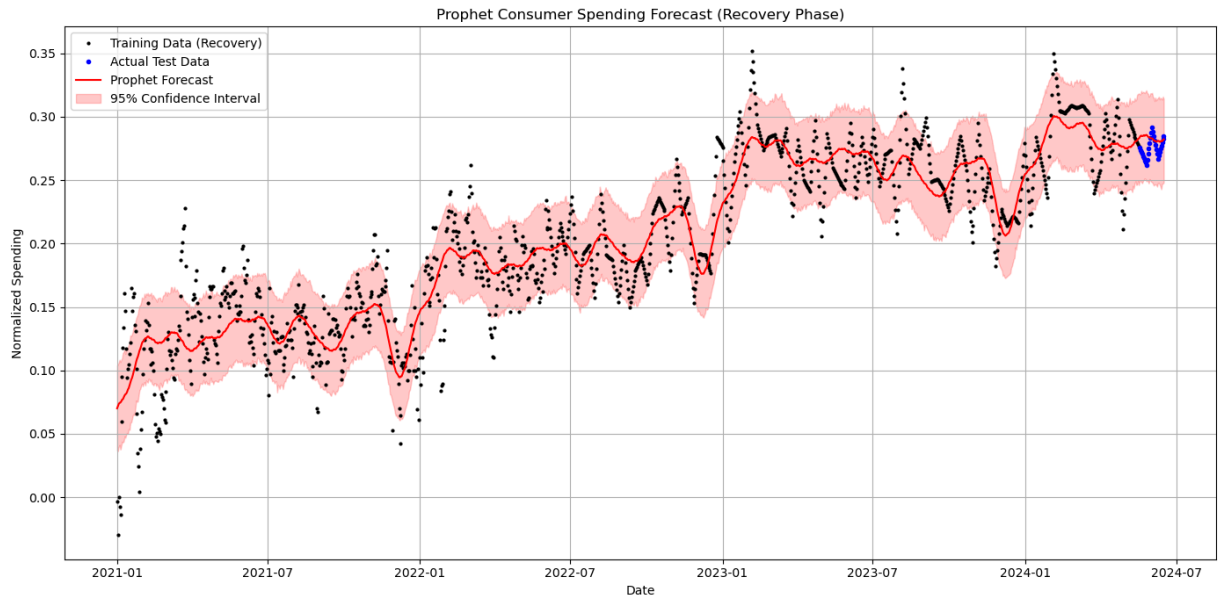
```
00:22:20 - cmdstanpy - INFO - Chain [1] done processing
```

--- Prophet Recovery Phase Evaluation ---

Mean Absolute Error (MAE): 0.009459

Mean Squared Error (MSE): 0.000132

Root Mean Squared Error (RMSE): 0.011510



MODEL 3: LSTM

```
In [12]: print("\n" + "="*80)
print("MODEL 3: LSTM - FULL DATASET")
print("="*80)

# Function to create sequences for LSTM
def create_sequences(data, seq_length):
    """Create sequences of seq_length days for LSTM training"""
    X, y = [], []
    for i in range(len(data) - seq_length):
        X.append(data[i:i + seq_length])
        y.append(data[i + seq_length])
    return np.array(X), np.array(y)

# Extract the data for modeling
data = df_timeseries['spend_all'].values.reshape(-1, 1)

# Normalize the data
scaler = MinMaxScaler(feature_range=(0, 1))
```

```

scaled_data = scaler.fit_transform(data)

# Define sequence length (window size)
sequence_length = 30 # 30 days of history to predict the next day

# Create sequences
X, y = create_sequences(scaled_data, sequence_length)

# Split into training and testing sets
X_train, X_test = X[:-test_size], X[-test_size:]
y_train, y_test = y[:-test_size], y[-test_size:]

print(f"Training data shape: X_train {X_train.shape}, y_train {y_train.shape}")
print(f"Testing data shape: X_test {X_test.shape}, y_test {y_test.shape}")

# Build LSTM model
print("Building and training LSTM model...")
lstm_model = Sequential([
    LSTM(units=50, return_sequences=True, input_shape=(sequence_length, 1)),
    Dropout(0.2),
    LSTM(units=50, return_sequences=False),
    Dropout(0.2),
    Dense(units=25),
    Dense(units=1)
])

# Compile the model
lstm_model.compile(optimizer='adam', loss='mean_squared_error')

# Early stopping
early_stop = EarlyStopping(monitor='val_loss', patience=10, verbose=1, restore_best

# Train the model
history = lstm_model.fit(
    X_train, y_train,
    epochs=100,
    batch_size=32,
    validation_split=0.2,
    callbacks=[early_stop],
    verbose=1
)

# Plot training history
plt.figure(figsize=(12, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('LSTM Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.savefig('lstm_training_loss.png')
plt.show()

# Make predictions on test data
lstm_predictions = lstm_model.predict(X_test)

```

```

# Inverse transform the predictions and actual values to original scale
lstm_predicted_values = scaler.inverse_transform(lstm_predictions)
lstm_actual_values = scaler.inverse_transform(y_test)

# Calculate evaluation metrics
mae_lstm = mean_absolute_error(lstm_actual_values, lstm_predicted_values)
mse_lstm = mean_squared_error(lstm_actual_values, lstm_predicted_values)
rmse_lstm = np.sqrt(mse_lstm)

print("\n--- LSTM Model Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_lstm:.6f}")
print(f"Mean Squared Error (MSE): {mse_lstm:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_lstm:.6f}")

# Add to results
add_result('LSTM', 'Full Data', mae_lstm, mse_lstm, rmse_lstm)

# Plot the predictions
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'].values[-len(lstm_actual_values):], lstm_actual_value)
plt.plot(df_timeseries['date'].values[-len(lstm_predicted_values):], lstm_predicted)
plt.title('LSTM Consumer Spending Forecast vs Actual')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.savefig('lstm_predictions.png')
plt.show()

# Forecast next 30 days
last_sequence = scaled_data[-sequence_length:]
next_30_days_scaled = []

# Iteratively predict each of the next 30 days
for _ in range(30):
    # Reshape the last sequence for prediction
    last_sequence_reshaped = last_sequence.reshape(1, sequence_length, 1)

    # Predict the next day
    next_day_scaled = lstm_model.predict(last_sequence_reshaped)

    # Append to our predictions
    next_30_days_scaled.append(next_day_scaled[0, 0])

    # Update the last sequence
    last_sequence = np.append(last_sequence[1:], next_day_scaled[0])
    last_sequence = last_sequence.reshape(-1, 1)

# Convert the predicted values back to the original scale
next_30_days_lstm = scaler.inverse_transform(np.array(next_30_days_scaled).reshape(
# Generate dates for the 30-day forecast
last_date = df_timeseries['date'].iloc[-1]
forecast_dates_lstm = pd.date_range(start=last_date + pd.Timedelta(days=1), periods

```

```

# Plot the forecast
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'].values[-90:], df_timeseries['spend_all'].values[-90:
plt.plot(forecast_dates_lstm, next_30_days_lstm, label='LSTM 30-Day Forecast', color=
plt.axvline(x=last_date, color='black', linestyle='--', label='Forecast Start')
plt.title('LSTM 30-Day Consumer Spending Forecast')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.savefig('lstm_future_forecast.png')
plt.show()

```

```

=====
MODEL 3: LSTM - FULL DATASET
=====

```

```

Training data shape: X_train (1935, 30, 1), y_train (1935, 1)

```

```

Testing data shape: X_test (30, 30, 1), y_test (30, 1)

```

```

Building and training LSTM model...

```

```


Epoch 1/100


```


```


D:\anaconda\Lib\site-packages\keras\src\layers\rnn\rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)


```


49/49  4s 20ms/step - loss: 0.0967 - val_loss: 0.0082
Epoch 2/100


49/49  1s 13ms/step - loss: 0.0070 - val_loss: 0.0013
Epoch 3/100


49/49  1s 13ms/step - loss: 0.0058 - val_loss: 0.0020
Epoch 4/100


49/49  1s 13ms/step - loss: 0.0047 - val_loss: 0.0022
Epoch 5/100


49/49  1s 13ms/step - loss: 0.0045 - val_loss: 0.0034
Epoch 6/100


49/49  1s 13ms/step - loss: 0.0043 - val_loss: 0.0022
Epoch 7/100


49/49  1s 13ms/step - loss: 0.0039 - val_loss: 0.0019
Epoch 8/100


49/49  1s 13ms/step - loss: 0.0043 - val_loss: 0.0014
Epoch 9/100


49/49  1s 13ms/step - loss: 0.0036 - val_loss: 0.0012
Epoch 10/100


49/49  1s 13ms/step - loss: 0.0033 - val_loss: 0.0014
Epoch 11/100


49/49  1s 13ms/step - loss: 0.0031 - val_loss: 0.0012
Epoch 12/100


49/49  1s 13ms/step - loss: 0.0033 - val_loss: 0.0011
Epoch 13/100


49/49  1s 13ms/step - loss: 0.0027 - val_loss: 0.0011
Epoch 14/100


49/49  1s 14ms/step - loss: 0.0029 - val_loss: 0.0015
Epoch 15/100


49/49  1s 13ms/step - loss: 0.0030 - val_loss: 0.0011
Epoch 16/100


49/49  1s 13ms/step - loss: 0.0026 - val_loss: 0.0011
Epoch 17/100


49/49  1s 14ms/step - loss: 0.0026 - val_loss: 0.0010
Epoch 18/100


49/49  1s 14ms/step - loss: 0.0022 - val_loss: 0.0012
Epoch 19/100


49/49  1s 14ms/step - loss: 0.0023 - val_loss: 0.0011
Epoch 20/100


49/49  1s 14ms/step - loss: 0.0022 - val_loss: 0.0014
Epoch 21/100


49/49  1s 13ms/step - loss: 0.0023 - val_loss: 0.0010
Epoch 22/100


49/49  1s 13ms/step - loss: 0.0020 - val_loss: 0.0013
Epoch 23/100


49/49  1s 13ms/step - loss: 0.0019 - val_loss: 9.7845e-04
Epoch 24/100


49/49  1s 13ms/step - loss: 0.0020 - val_loss: 9.2242e-04
Epoch 25/100


49/49  1s 14ms/step - loss: 0.0019 - val_loss: 0.0012
Epoch 26/100


49/49  1s 15ms/step - loss: 0.0018 - val_loss: 8.9818e-04
Epoch 27/100


49/49  1s 14ms/step - loss: 0.0016 - val_loss: 9.9338e-04
Epoch 28/100


49/49  1s 14ms/step - loss: 0.0017 - val_loss: 7.6063e-04
Epoch 29/100


49/49  1s 14ms/step - loss: 0.0015 - val_loss: 6.4074e-04
Epoch 30/100


49/49  1s 15ms/step - loss: 0.0014 - val_loss: 4.8505e-04
Epoch 31/100


49/49  1s 14ms/step - loss: 0.0013 - val_loss: 6.4297e-04
Epoch 32/100


49/49  1s 13ms/step - loss: 0.0013 - val_loss: 4.1550e-04
Epoch 33/100


49/49  1s 13ms/step - loss: 0.0012 - val_loss: 3.2900e-04
Epoch 34/100


49/49  1s 13ms/step - loss: 0.0012 - val_loss: 3.7441e-04
Epoch 35/100


49/49  1s 13ms/step - loss: 0.0011 - val_loss: 3.4294e-04
Epoch 36/100


49/49  1s 13ms/step - loss: 0.0011 - val_loss: 2.3725e-04
Epoch 37/100


49/49  1s 13ms/step - loss: 9.9127e-04 - val_loss: 3.5598e-04
Epoch 38/100


49/49  1s 13ms/step - loss: 0.0011 - val_loss: 2.6694e-04
Epoch 39/100


49/49  1s 13ms/step - loss: 0.0010 - val_loss: 1.9393e-04
Epoch 40/100


49/49  1s 13ms/step - loss: 8.9301e-04 - val_loss: 3.9031e-04
Epoch 41/100


49/49  1s 13ms/step - loss: 8.8816e-04 - val_loss: 2.1568e-04
Epoch 42/100


49/49  1s 13ms/step - loss: 9.3442e-04 - val_loss: 2.2907e-04
Epoch 43/100


49/49  1s 13ms/step - loss: 9.3460e-04 - val_loss: 3.4992e-04
Epoch 44/100


49/49  1s 14ms/step - loss: 8.7873e-04 - val_loss: 5.4185e-04
Epoch 45/100


49/49  1s 13ms/step - loss: 8.1239e-04 - val_loss: 2.3649e-04
Epoch 46/100


49/49  1s 14ms/step - loss: 8.9229e-04 - val_loss: 1.7837e-04
Epoch 47/100


49/49  1s 13ms/step - loss: 8.3167e-04 - val_loss: 2.1110e-04
Epoch 48/100


49/49  1s 13ms/step - loss: 8.0729e-04 - val_loss: 2.1540e-04
Epoch 49/100


49/49  1s 13ms/step - loss: 8.0825e-04 - val_loss: 2.2702e-04
Epoch 50/100


49/49  1s 13ms/step - loss: 8.6381e-04 - val_loss: 4.5304e-04
Epoch 51/100


49/49  1s 13ms/step - loss: 8.2115e-04 - val_loss: 2.8333e-04
Epoch 52/100

49/49  1s 13ms/step - loss: 7.4998e-04 - val_loss: 3.5312e-04
Epoch 53/100

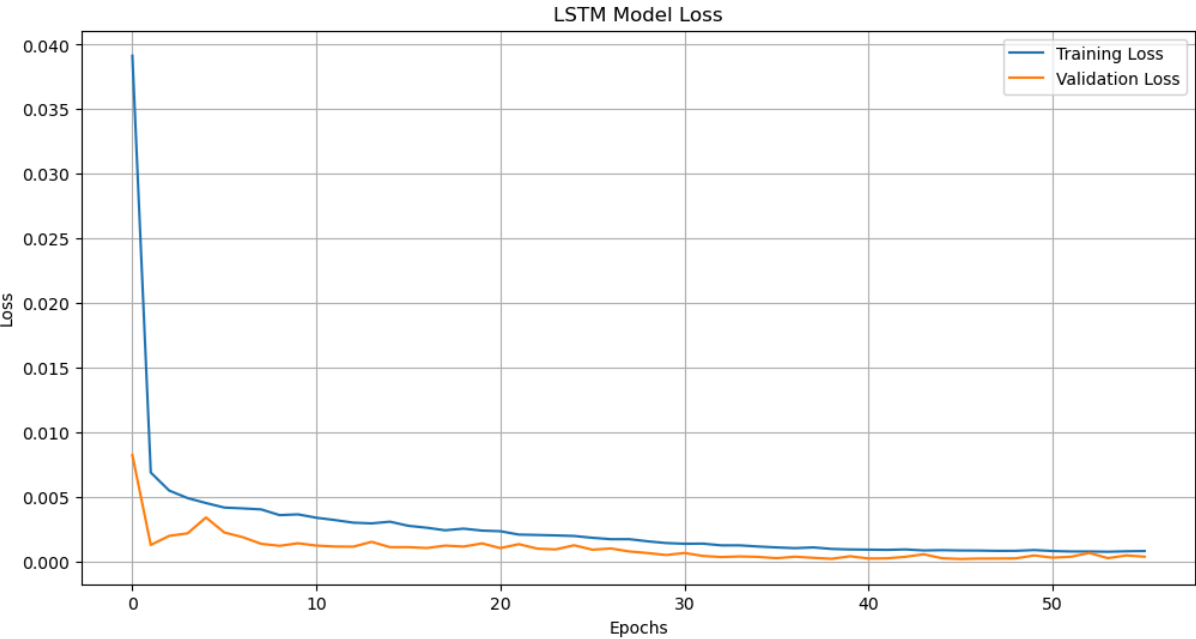
49/49  1s 13ms/step - loss: 7.4780e-04 - val_loss: 6.4702e-04
Epoch 54/100

49/49  1s 14ms/step - loss: 7.8569e-04 - val_loss: 2.4506e-04
Epoch 55/100

49/49  1s 13ms/step - loss: 7.6455e-04 - val_loss: 4.5197e-04
Epoch 56/100

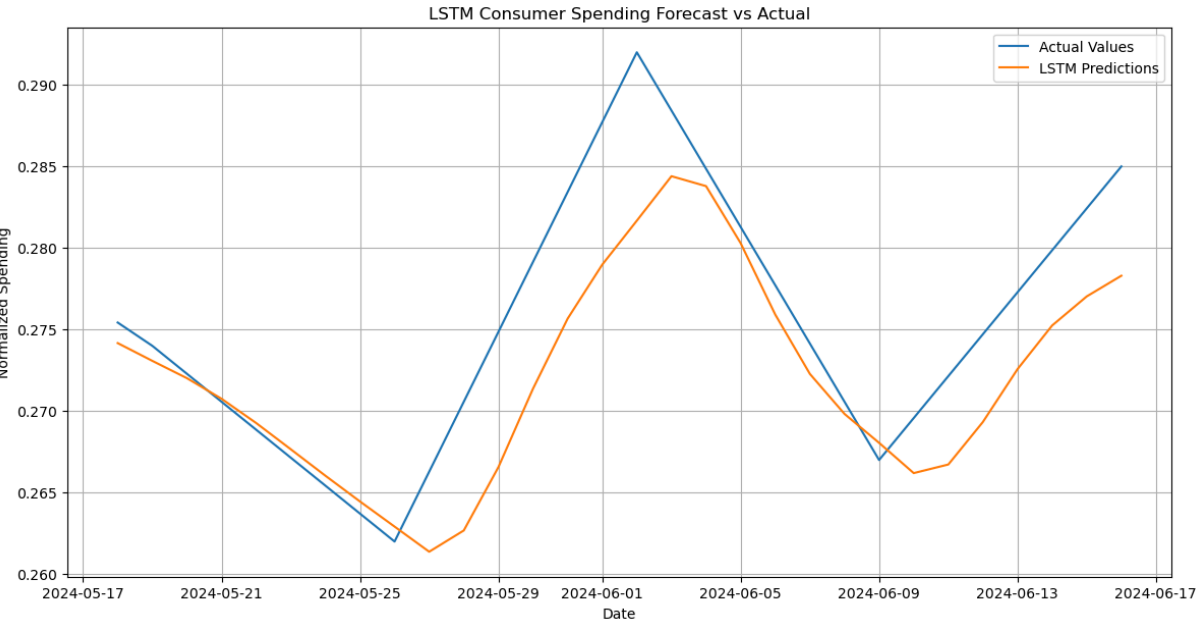
49/49  1s 13ms/step - loss: 7.8030e-04 - val_loss: 3.5754e-04

Epoch 56: early stopping
Restoring model weights from the end of the best epoch: 46.

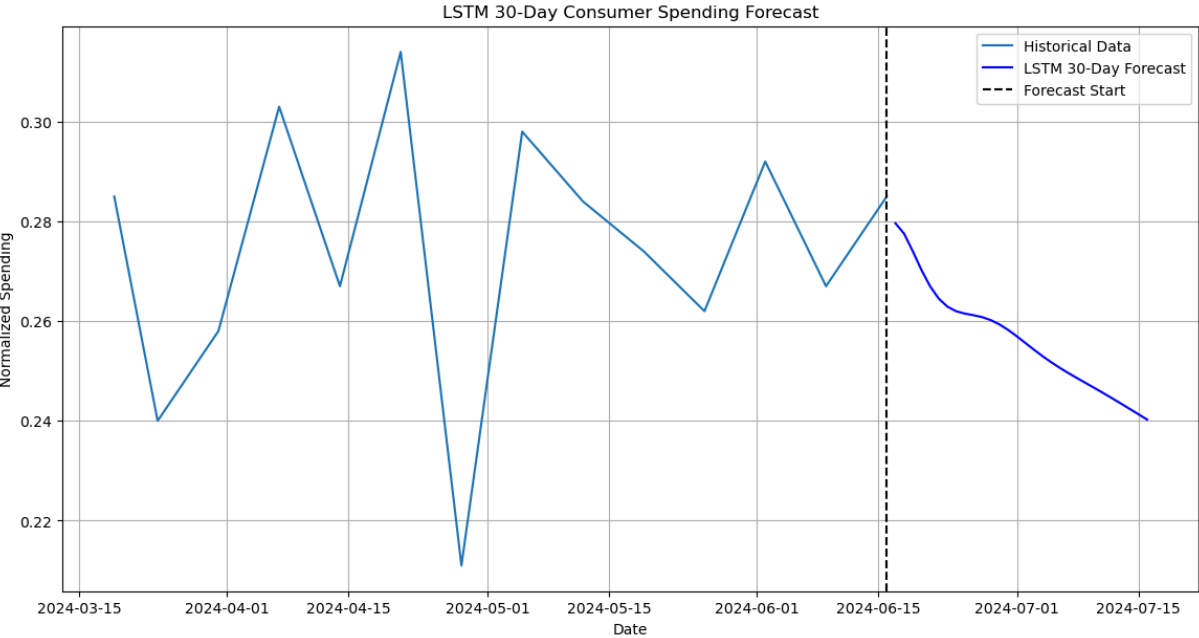


1/1 ————— 0s 269ms/step

--- LSTM Model Evaluation ---
Mean Absolute Error (MAE): 0.003622
Mean Squared Error (MSE): 0.000023
Root Mean Squared Error (RMSE): 0.004752



1/1	0s	277ms/step
1/1	0s	17ms/step
1/1	0s	21ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	15ms/step
1/1	0s	16ms/step
1/1	0s	17ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	16ms/step
1/1	0s	15ms/step
1/1	0s	15ms/step
1/1	0s	15ms/step
1/1	0s	21ms/step
1/1	0s	16ms/step
1/1	0s	19ms/step
1/1	0s	18ms/step
1/1	0s	18ms/step
1/1	0s	16ms/step
1/1	0s	15ms/step
1/1	0s	17ms/step
1/1	0s	17ms/step



Run LSTM on recovery phase data

```
-----  
-----  
  
In [14]: print("\n" + "="*80)  
print("MODEL 3: LSTM - RECOVERY PHASE")  
print("="*80)  
  
# Extract recovery phase data  
recovery_data_array = recovery_data['spend_all'].values.reshape(-1, 1)  
  
# Normalize the recovery data  
scaler_recovery = MinMaxScaler(feature_range=(0, 1))  
scaled_recovery_data = scaler_recovery.fit_transform(recovery_data_array)  
  
# Create sequences for recovery data  
X_recovery, y_recovery = create_sequences(scaled_recovery_data, sequence_length)  
  
# Split into training and testing  
X_recovery_train, X_recovery_test = X_recovery[:-test_size], X_recovery[-test_size:  
y_recovery_train, y_recovery_test = y_recovery[:-test_size], y_recovery[-test_size:  
  
# Build LSTM model for recovery data  
lstm_model_recovery = Sequential([  
    LSTM(units=50, return_sequences=True, input_shape=(sequence_length, 1)),  
    Dropout(0.2),  
    LSTM(units=50, return_sequences=False),  
    Dropout(0.2),  
    Dense(units=25),  
    Dense(units=1)  
)  
  
# Compile the model  
lstm_model_recovery.compile(optimizer='adam', loss='mean_squared_error')  
  
# Train the model on recovery data  
lstm_model_recovery.fit(  
    X_recovery_train, y_recovery_train,  
    epochs=100,  
    batch_size=32,  
    validation_split=0.2,  
    callbacks=[early_stop],  
    verbose=1  
)  
  
# Make predictions on test data  
lstm_predictions_recovery = lstm_model_recovery.predict(X_recovery_test)  
  
# Inverse transform the predictions and actual values  
lstm_predicted_values_recovery = scaler_recovery.inverse_transform(lstm_predictions_recovery)  
lstm_actual_values_recovery = scaler_recovery.inverse_transform(y_recovery_test)
```

```

# Calculate evaluation metrics
mae_lstm_recovery = mean_absolute_error(lstm_actual_values_recovery, lstm_predicted_values_recovery)
mse_lstm_recovery = mean_squared_error(lstm_actual_values_recovery, lstm_predicted_values_recovery)
rmse_lstm_recovery = np.sqrt(mse_lstm_recovery)

print("\n--- LSTM Recovery Phase Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_lstm_recovery:.6f}")
print(f"Mean Squared Error (MSE): {mse_lstm_recovery:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_lstm_recovery:.6f}")

# Add to results
add_result('LSTM', 'Recovery Phase', mae_lstm_recovery, mse_lstm_recovery, rmse_lstm_recovery)

# Plot recovery results
plt.figure(figsize=(14, 7))
plt.plot(recovery_data['date'].values[-len(lstm_actual_values_recovery):], lstm_actual_values_recovery, label='Actual')
plt.plot(recovery_data['date'].values[-len(lstm_predicted_values_recovery):], lstm_predicted_values_recovery, label='Predicted')
plt.title('LSTM Consumer Spending Forecast vs Actual (Recovery Phase)')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.savefig('lstm_predictions_recovery.png')
plt.show()

```

```

=====
MODEL 3: LSTM - RECOVERY PHASE
=====

```

```
Epoch 1/100
```

```

D:\anaconda\Lib\site-packages\keras\src\layers\rnn\rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(**kwargs)

```

31/31 4s 25ms/step - loss: 0.0934 - val_loss: 0.0042
Epoch 2/100

31/31 0s 13ms/step - loss: 0.0114 - val_loss: 0.0112
Epoch 3/100

31/31 0s 14ms/step - loss: 0.0093 - val_loss: 0.0089
Epoch 4/100

31/31 0s 14ms/step - loss: 0.0080 - val_loss: 0.0107
Epoch 5/100

31/31 0s 15ms/step - loss: 0.0078 - val_loss: 0.0078
Epoch 6/100

31/31 0s 14ms/step - loss: 0.0076 - val_loss: 0.0058
Epoch 7/100

31/31 0s 14ms/step - loss: 0.0074 - val_loss: 0.0050
Epoch 8/100

31/31 0s 13ms/step - loss: 0.0064 - val_loss: 0.0045
Epoch 9/100

31/31 0s 14ms/step - loss: 0.0069 - val_loss: 0.0044
Epoch 10/100

31/31 0s 13ms/step - loss: 0.0064 - val_loss: 0.0052
Epoch 10: early stopping
Restoring model weights from the end of the best epoch: 1.

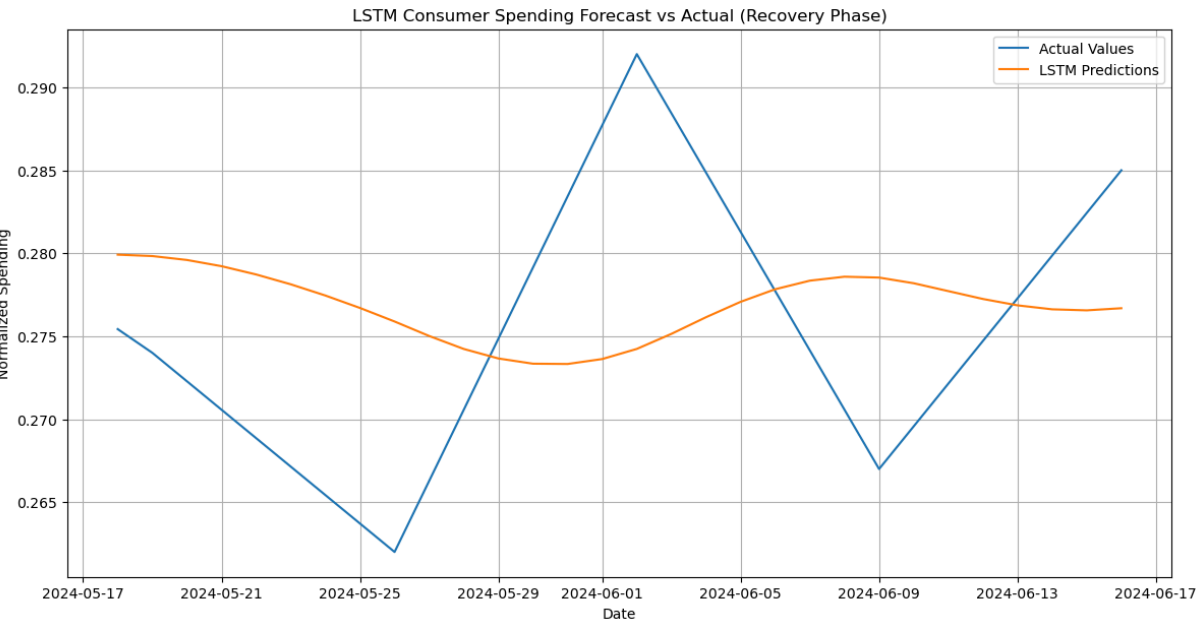
1/1 0s 332ms/step

--- LSTM Recovery Phase Evaluation ---

Mean Absolute Error (MAE): 0.007734

Mean Squared Error (MSE): 0.000079

Root Mean Squared Error (RMSE): 0.008866



MODEL 4: TRANSFORMER

```

-----

-----

In [23]: print("\n" + "="*80)
print("MODEL 4: TRANSFORMER - FULL DATASET")
print("="*80)

# Function for positional encoding
def positional_encoding(sequence_length, d_model):
    """Generate positional encoding for Transformer model"""
    positions = np.arange(sequence_length)[: , np.newaxis]
    angles = np.arange(d_model)[np.newaxis, :] / np.power(10000, 2 * (np.arange(d_m

    # Apply sin to even indices in the array
    sines = np.sin(positions * angles[:, 0::2])
    # Apply cos to odd indices in the array
    cosines = np.cos(positions * angles[:, 1::2])

    # Combine sin and cos positional encodings
    pos_encoding = np.zeros((sequence_length, d_model))
    pos_encoding[:, 0::2] = sines
    pos_encoding[:, 1::2] = cosines

    return tf.cast(pos_encoding, dtype=tf.float32)

# Create Transformer blocks
def transformer_block(inputs, d_model, num_heads, ff_dim, dropout=0.1):
    """Transformer block with multi-head attention"""
    # Multi-head self-attention
    attention_output = MultiHeadAttention(
        num_heads=num_heads, key_dim=d_model // num_heads
    )(inputs, inputs)
    attention_output = Dropout(dropout)(attention_output)
    attention_output = LayerNormalization(epsilon=1e-6)(inputs + attention_output)

    # Feed forward network
    ff_output = Dense(ff_dim, activation="relu")(attention_output)
    ff_output = Dense(d_model)(ff_output)
    ff_output = Dropout(dropout)(ff_output)
    return LayerNormalization(epsilon=1e-6)(attention_output + ff_output)

# Build the Transformer model
print("Building Transformer model...")

# Define model parameters
d_model = 32 # Embedding dimension
num_heads = 4 # Number of attention heads
ff_dim = 64 # Feed forward network dimension
dropout_rate = 0.1

# Input Layer
inputs = Input(shape=(sequence_length, 1))

# Embedding Layer (expand 1D data to d_model dimensions)

```

```
x = Dense(d_model)(inputs)

# Add positional encoding
pos_encoding = positional_encoding(sequence_length, d_model)
x = x + pos_encoding

# Transformer blocks
x = transformer_block(x, d_model, num_heads, ff_dim, dropout_rate)
x = transformer_block(x, d_model, num_heads, ff_dim, dropout_rate)

# Global pooling
x = GlobalAveragePooling1D()(x)

# Output layer
outputs = Dense(1)(x)

# Create and compile model
transformer_model = Model(inputs=inputs, outputs=outputs)
transformer_model.compile(optimizer='adam', loss='mean_squared_error')

# Train model
print("Training Transformer model...")
transformer_history = transformer_model.fit(
    X_train, y_train,
    epochs=100,
    batch_size=32,
    validation_split=0.2,
    callbacks=[early_stop],
    verbose=1)
```

=====

MODEL 4: TRANSFORMER - FULL DATASET

=====

Building Transformer model...

Training Transformer model...

Epoch 1/100

49/49 ————— 8s 16ms/step - loss: 0.0575 - val_loss: 0.0026

Epoch 2/100

49/49 ————— 0s 10ms/step - loss: 0.0066 - val_loss: 0.0036

Epoch 3/100

49/49 ————— 0s 9ms/step - loss: 0.0061 - val_loss: 0.0053

Epoch 4/100

49/49 ————— 0s 9ms/step - loss: 0.0064 - val_loss: 0.0032

Epoch 5/100

49/49 ————— 0s 9ms/step - loss: 0.0050 - val_loss: 0.0049

Epoch 6/100

49/49 ————— 0s 9ms/step - loss: 0.0049 - val_loss: 0.0045

Epoch 7/100

49/49 ————— 0s 9ms/step - loss: 0.0044 - val_loss: 0.0033

Epoch 8/100

49/49 ————— 0s 10ms/step - loss: 0.0043 - val_loss: 0.0041

Epoch 9/100

49/49 ————— 0s 10ms/step - loss: 0.0045 - val_loss: 0.0036

Epoch 10/100

49/49 ————— 0s 9ms/step - loss: 0.0046 - val_loss: 0.0017

Epoch 10: early stopping

Restoring model weights from the end of the best epoch: 1.

```
In [24]: # Plot training history for transformer
plt.figure(figsize=(12, 6))
plt.plot(transformer_history.history['loss'], label='Training Loss')
plt.plot(transformer_history.history['val_loss'], label='Validation Loss')
plt.title('Transformer Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.savefig('transformer_training_loss.png')
plt.show()

# Make predictions on test data
transformer_predictions = transformer_model.predict(X_test)

# Inverse transform the predictions and actual values
transformer_predicted_values = scaler.inverse_transform(transformer_predictions)
transformer_actual_values = scaler.inverse_transform(y_test)

# Calculate evaluation metrics
mae_transformer = mean_absolute_error(transformer_actual_values, transformer_predict
mse_transformer = mean_squared_error(transformer_actual_values, transformer_predict
rmse_transformer = np.sqrt(mse_transformer)

print("\n--- Transformer Model Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_transformer:.6f}")
print(f"Mean Squared Error (MSE): {mse_transformer:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_transformer:.6f}")
```



```

# Add to results
add_result('Transformer', 'Full Data', mae_transformer, mse_transformer, rmse_trans

# Plot the predictions
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'].values[-len(transformer_actual_values):], transformer
plt.plot(df_timeseries['date'].values[-len(transformer_predicted_values):], transfo
plt.title('Transformer Consumer Spending Forecast vs Actual')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.savefig('transformer_predictions.png')
plt.show()

# Forecast next 30 days with transformer model
last_sequence = scaled_data[-sequence_length:]
transformer_next_30_days_scaled = []

# Iteratively predict each of the next 30 days
for _ in range(30):
    # Reshape the last sequence for prediction
    last_sequence_reshaped = last_sequence.reshape(1, sequence_length, 1)

    # Predict the next day
    next_day_scaled = transformer_model.predict(last_sequence_reshaped)

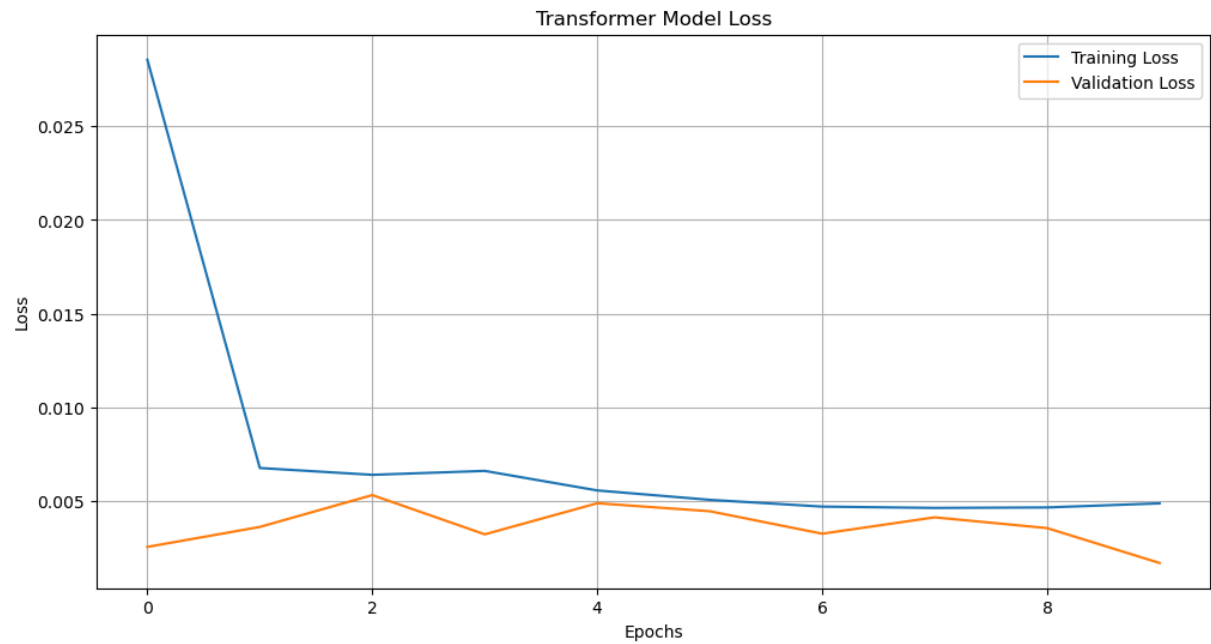
    # Append to our predictions
    transformer_next_30_days_scaled.append(next_day_scaled[0, 0])

    # Update the last sequence
    last_sequence = np.append(last_sequence[1:], next_day_scaled[0])
    last_sequence = last_sequence.reshape(-1, 1)

# Convert the predicted values back to the original scale
next_30_days_transformer = scaler.inverse_transform(np.array(transformer_next_30_da

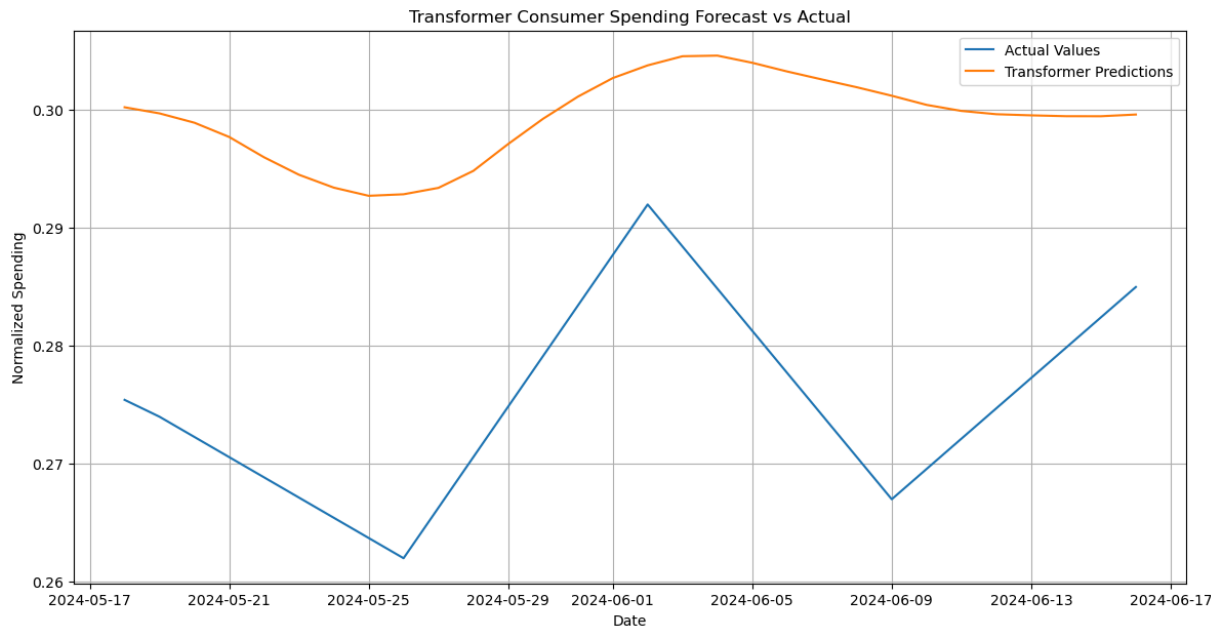
# Plot the forecast
plt.figure(figsize=(14, 7))
plt.plot(df_timeseries['date'].values[-90:], df_timeseries['spend_all'].values[-90:
plt.plot(forecast_dates_lstm, next_30_days_transformer, label='Transformer 30-Day F
plt.axvline(x=last_date, color='black', linestyle='--', label='Forecast Start')
plt.title('Transformer 30-Day Consumer Spending Forecast')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.savefig('transformer_future_forecast.png')
plt.show()

```

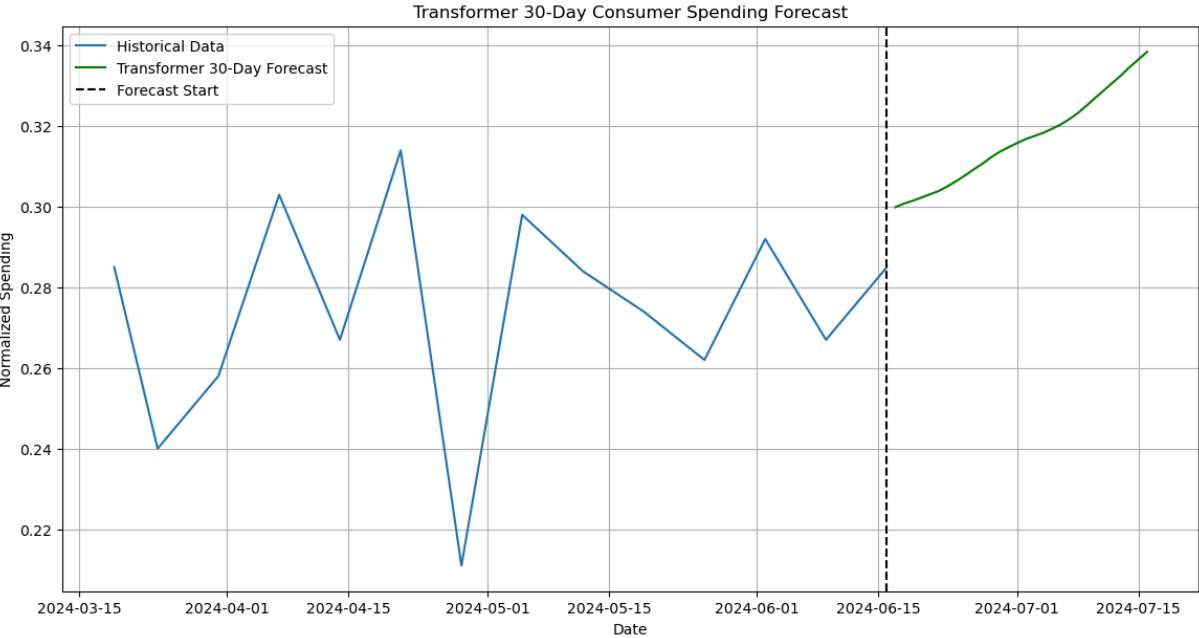


1/1 0s 270ms/step

--- Transformer Model Evaluation ---
Mean Absolute Error (MAE): 0.024015
Mean Squared Error (MSE): 0.000607
Root Mean Squared Error (RMSE): 0.024639



1/1		0s	272ms/step
1/1		0s	18ms/step
1/1		0s	17ms/step
1/1		0s	17ms/step
1/1		0s	17ms/step
1/1		0s	16ms/step
1/1		0s	18ms/step
1/1		0s	15ms/step
1/1		0s	17ms/step
1/1		0s	18ms/step
1/1		0s	17ms/step
1/1		0s	17ms/step
1/1		0s	16ms/step
1/1		0s	17ms/step
1/1		0s	16ms/step
1/1		0s	16ms/step
1/1		0s	16ms/step
1/1		0s	16ms/step
1/1		0s	17ms/step
1/1		0s	17ms/step
1/1		0s	15ms/step
1/1		0s	19ms/step
1/1		0s	17ms/step
1/1		0s	15ms/step
1/1		0s	17ms/step
1/1		0s	15ms/step
1/1		0s	16ms/step
1/1		0s	18ms/step
1/1		0s	15ms/step
1/1		0s	16ms/step



Run Transformer on recovery phase data

```

In [26]: print("\n" + "="*80)
print("MODEL 4: TRANSFORMER - RECOVERY PHASE")
print("="*80)

# Build a new transformer model for recovery data
transformer_model_recovery = Model(inputs=inputs, outputs=outputs)
transformer_model_recovery.compile(optimizer='adam', loss='mean_squared_error')

# Train on recovery data
transformer_model_recovery.fit(
    X_recovery_train, y_recovery_train,
    epochs=100,
    batch_size=32,
    validation_split=0.2,
    callbacks=[early_stop],
    verbose=1
)

# Make predictions on test data
transformer_predictions_recovery = transformer_model_recovery.predict(X_recovery_te

# Inverse transform the predictions and actual values
transformer_predicted_values_recovery = scaler_recovery.inverse_transform(transformer_predictions_recovery)
transformer_actual_values_recovery = scaler_recovery.inverse_transform(y_recovery_train)

# Calculate evaluation metrics
mae_transformer_recovery = mean_absolute_error(transformer_actual_values_recovery, transformer_predictions_recovery)
mse_transformer_recovery = mean_squared_error(transformer_actual_values_recovery, transformer_predictions_recovery)
rmse_transformer_recovery = np.sqrt(mse_transformer_recovery)

print("\n--- Transformer Recovery Phase Evaluation ---")
print(f"Mean Absolute Error (MAE): {mae_transformer_recovery:.6f}")
print(f"Mean Squared Error (MSE): {mse_transformer_recovery:.6f}")
print(f"Root Mean Squared Error (RMSE): {rmse_transformer_recovery:.6f}")

# Add to results
add_result('Transformer', 'Recovery Phase', mae_transformer_recovery, mse_transformer_recovery, rmse_transformer_recovery)

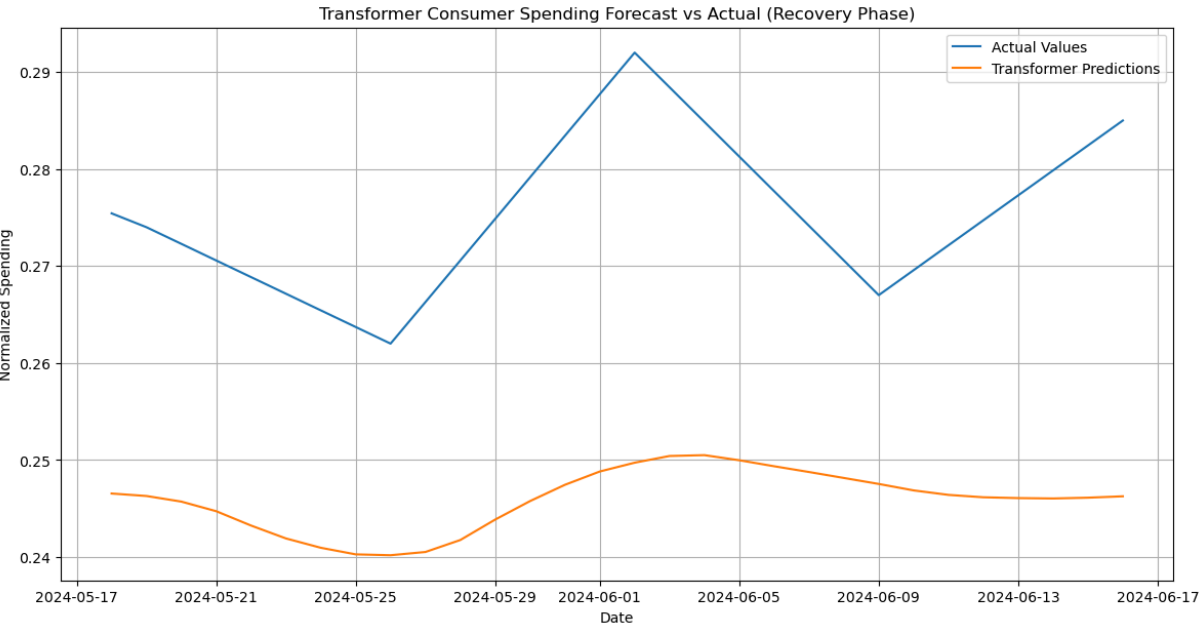
# Plot recovery results
plt.figure(figsize=(14, 7))
plt.plot(recovery_data['date'].values[-len(transformer_actual_values_recovery):], transformer_actual_values_recovery, label='Actual')
plt.plot(recovery_data['date'].values[-len(transformer_predictions_recovery):], transformer_predictions_recovery, label='Predicted')
plt.title('Transformer Consumer Spending Forecast vs Actual (Recovery Phase)')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)

```

```
plt.savefig('transformer_predictions_recovery.png')
plt.show()
```

```
=====
MODEL 4: TRANSFORMER - RECOVERY PHASE
=====
Epoch 1/100
31/31 ----- 8s 21ms/step - loss: 0.0754 - val_loss: 0.0102
Epoch 2/100
31/31 ----- 0s 9ms/step - loss: 0.0111 - val_loss: 0.0048
Epoch 3/100
31/31 ----- 0s 9ms/step - loss: 0.0076 - val_loss: 0.0073
Epoch 4/100
31/31 ----- 0s 9ms/step - loss: 0.0088 - val_loss: 0.0051
Epoch 5/100
31/31 ----- 0s 9ms/step - loss: 0.0079 - val_loss: 0.0067
Epoch 6/100
31/31 ----- 0s 9ms/step - loss: 0.0087 - val_loss: 0.0075
Epoch 7/100
31/31 ----- 0s 9ms/step - loss: 0.0091 - val_loss: 0.0082
Epoch 8/100
31/31 ----- 0s 9ms/step - loss: 0.0084 - val_loss: 0.0063
Epoch 9/100
31/31 ----- 0s 9ms/step - loss: 0.0078 - val_loss: 0.0055
Epoch 10/100
31/31 ----- 0s 9ms/step - loss: 0.0074 - val_loss: 0.0042
Epoch 10: early stopping
Restoring model weights from the end of the best epoch: 1.
1/1 ----- 0s 260ms/step
```

```
--- Transformer Recovery Phase Evaluation ---
Mean Absolute Error (MAE): 0.029401
Mean Squared Error (MSE): 0.000897
Root Mean Squared Error (RMSE): 0.029956
```



MODEL COMPARISON AND FINAL ANALYSIS

```
In [28]: print("\n" + "="*80)
print("MODEL COMPARISON AND FINAL ANALYSIS")
print("="*80)

# Display the comparison table
print("\nModel Performance Comparison:")
print(results_df)

# Sort by RMSE for ranking
ranked_results = results_df.sort_values(by='RMSE')
print("\nModels Ranked by Performance (RMSE):")
print(ranked_results)

# Create a function to plot comparison charts
def plot_metric_comparison(metric):
    plt.figure(figsize=(12, 6))

    # Create grouped bar chart
    sns.barplot(x='Model', y=metric, hue='Data Type', data=results_df)

    plt.title(f'Model Comparison by {metric}')
    plt.ylabel(metric)
    plt.grid(True, axis='y')
    plt.tight_layout()
    plt.savefig(f'comparison_{metric}.png')
    plt.show()

# Plot comparisons for each metric
plot_metric_comparison('MAE')
plot_metric_comparison('MSE')
plot_metric_comparison('RMSE')

# Combined metric visualization
plt.figure(figsize=(15, 10))

# RMSE subplot
plt.subplot(3, 1, 1)
sns.barplot(x='Model', y='RMSE', hue='Data Type', data=results_df)
plt.title('Root Mean Squared Error (RMSE) Comparison')
```

```

plt.grid(True, axis='y')

# MAE subplot
plt.subplot(3, 1, 2)
sns.barplot(x='Model', y='MAE', hue='Data Type', data=results_df)
plt.title('Mean Absolute Error (MAE) Comparison')
plt.grid(True, axis='y')

# MSE subplot
plt.subplot(3, 1, 3)
sns.barplot(x='Model', y='MSE', hue='Data Type', data=results_df)
plt.title('Mean Squared Error (MSE) Comparison')
plt.grid(True, axis='y')

plt.tight_layout()
plt.savefig('combined_metrics_comparison.png')
plt.show()

# Plot all model future forecasts on one chart
plt.figure(figsize=(14, 7))
# Plot historical data
plt.plot(df_timeseries['date'].values[-90:], df_timeseries['spend_all'].values[-90:]

# Plot each model's forecast
plt.plot(future_dates, future_forecast_arima, 'r-', label='ARIMA Forecast')
plt.plot(future_forecast['ds'][-30:], future_forecast['yhat'][-30:], 'b-', label='P
plt.plot(forecast_dates_lstm, next_30_days_lstm, 'g-', label='LSTM Forecast')
plt.plot(forecast_dates_lstm, next_30_days_transformer, 'm-', label='Transformer Fo

# Add vertical line at forecast start
plt.axvline(x=last_date, color='black', linestyle='--', label='Forecast Start')

plt.title('30-Day Consumer Spending Forecast Comparison')
plt.xlabel('Date')
plt.ylabel('Normalized Spending')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.savefig('all_models_forecast_comparison.png')
plt.show()

```

=====

MODEL COMPARISON AND FINAL ANALYSIS

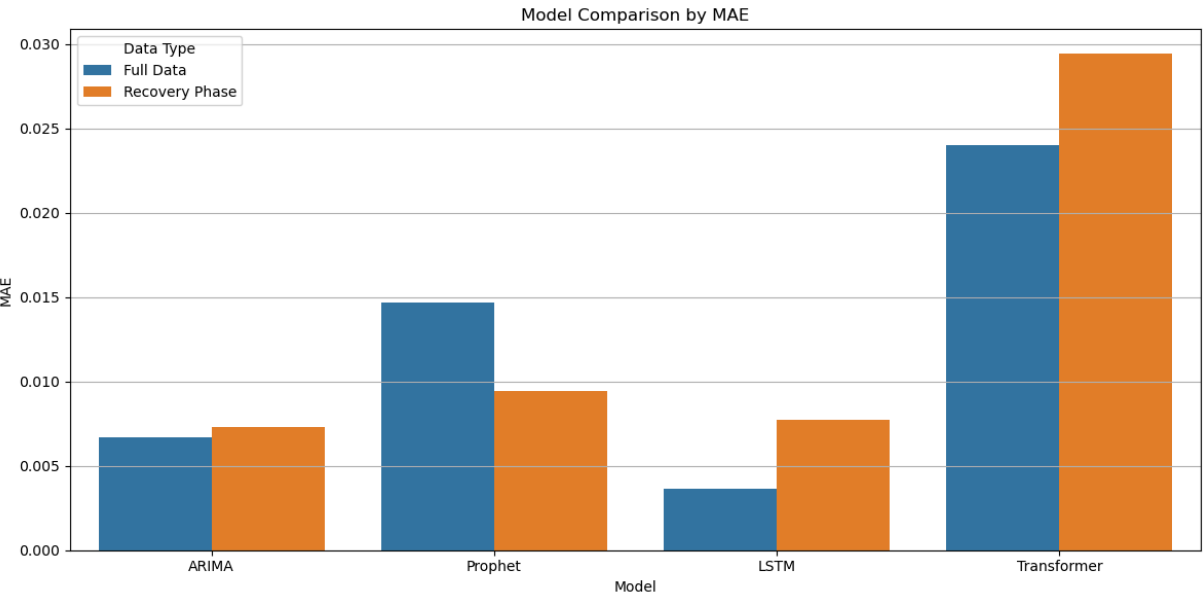
=====

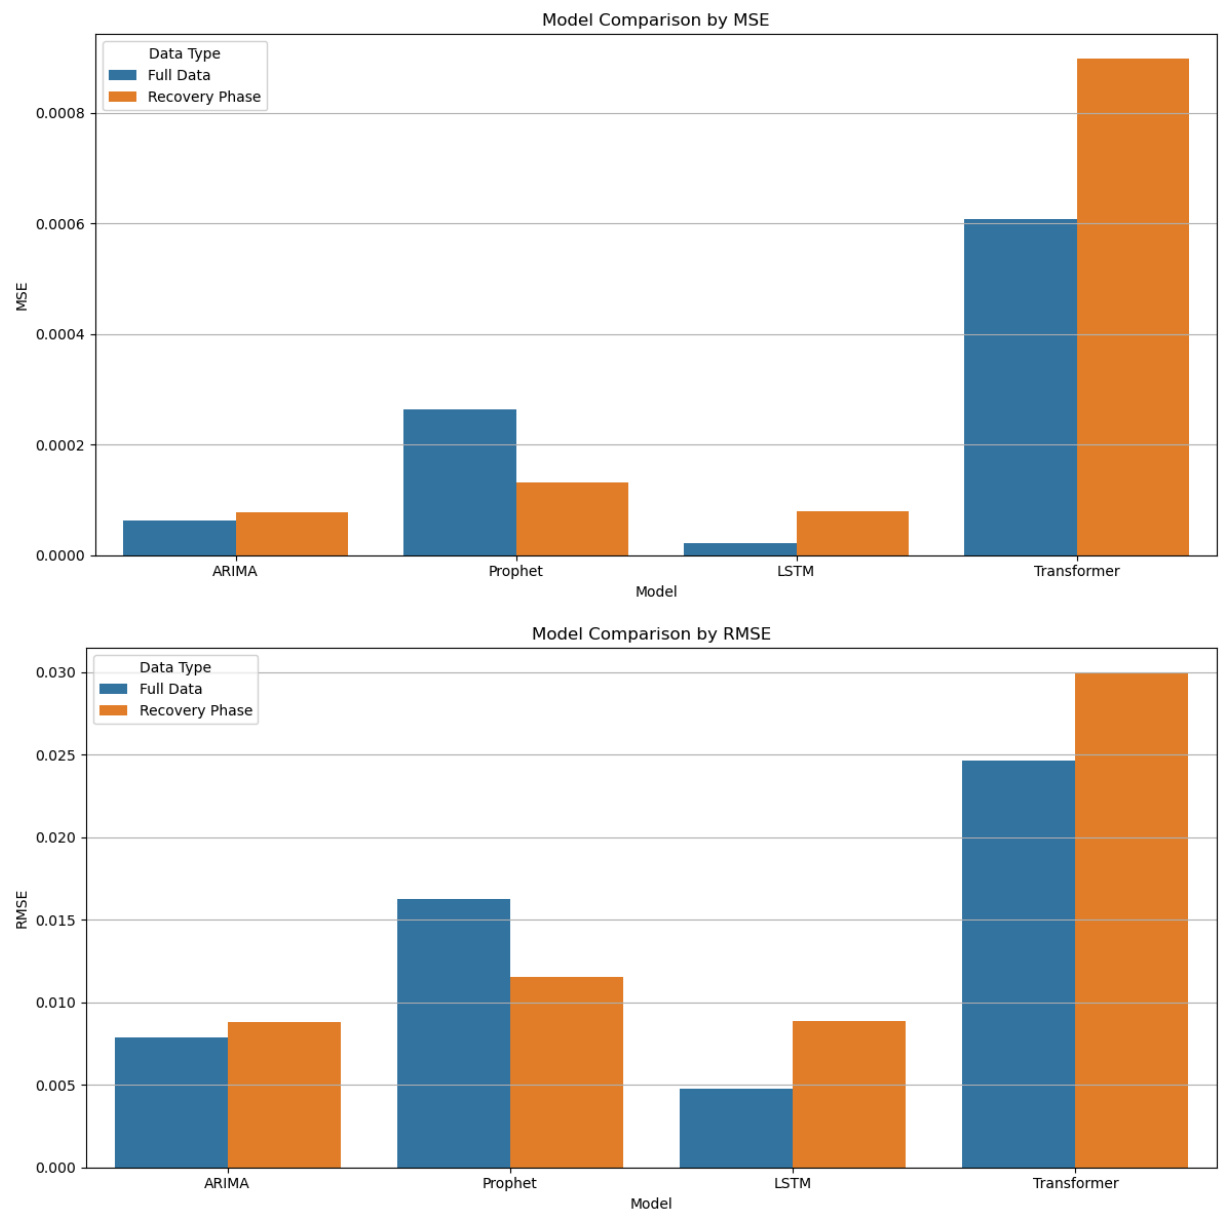
Model Performance Comparison:

	Model	Data Type	MAE	MSE	RMSE
0	ARIMA	Full Data	0.006686	0.000062	0.007886
1	ARIMA	Recovery Phase	0.007297	0.000078	0.008813
2	Prophet	Full Data	0.014657	0.000264	0.016253
3	Prophet	Recovery Phase	0.009459	0.000132	0.011510
4	LSTM	Full Data	0.003622	0.000023	0.004752
5	LSTM	Recovery Phase	0.007734	0.000079	0.008866
6	Transformer	Full Data	0.024015	0.000607	0.024639
7	Transformer	Recovery Phase	0.029401	0.000897	0.029956

Models Ranked by Performance (RMSE):

	Model	Data Type	MAE	MSE	RMSE
4	LSTM	Full Data	0.003622	0.000023	0.004752
0	ARIMA	Full Data	0.006686	0.000062	0.007886
1	ARIMA	Recovery Phase	0.007297	0.000078	0.008813
5	LSTM	Recovery Phase	0.007734	0.000079	0.008866
3	Prophet	Recovery Phase	0.009459	0.000132	0.011510
2	Prophet	Full Data	0.014657	0.000264	0.016253
6	Transformer	Full Data	0.024015	0.000607	0.024639
7	Transformer	Recovery Phase	0.029401	0.000897	0.029956







FINAL SUMMARY AND INTERPRETATION

```

In [31]: print("\n" + "="*80)
print("FINAL SUMMARY AND INTERPRETATION")
print("="*80)

# Get the best model for full data
best_full = ranked_results[ranked_results['Data Type'] == 'Full Data'].iloc[0]
# Get the best model for recovery phase
best_recovery = ranked_results[ranked_results['Data Type'] == 'Recovery Phase'].iloc[0]

print("\nBest Model for Full Dataset:")
print(f"Model: {best_full['Model']}")
print(f"RMSE: {best_full['RMSE']:.6f}")
print(f"MAE: {best_full['MAE']:.6f}")

print("\nBest Model for Recovery Phase Dataset:")
print(f"Model: {best_recovery['Model']}")
print(f"RMSE: {best_recovery['RMSE']:.6f}")
print(f"MAE: {best_recovery['MAE']:.6f}")

# Analysis of results
print("\nKey Findings and Interpretation:")

# Determine which dataset yielded better results
full_mean_rmse = results_df[results_df['Data Type'] == 'Full Data']['RMSE'].mean()
recovery_mean_rmse = results_df[results_df['Data Type'] == 'Recovery Phase']['RMSE'].mean()

if full_mean_rmse < recovery_mean_rmse:
    better_data = "full historical dataset"
    improvement = ((recovery_mean_rmse - full_mean_rmse) / recovery_mean_rmse) * 100
else:
    better_data = "recovery phase dataset"
    improvement = ((full_mean_rmse - recovery_mean_rmse) / full_mean_rmse) * 100

print(f"1. The {better_data} generally produced better forecasts across models (by")
print(f"2. The {best_full['Model']} model performed best on the full dataset with RMSE of {best_full['RMSE']:.6f}")
print(f"3. The {best_recovery['Model']} model performed best on the recovery phase with RMSE of {best_recovery['RMSE']:.6f}")

# Compare traditional vs deep Learning approaches
traditional_models = ['ARIMA', 'Prophet']
dl_models = ['LSTM', 'Transformer']

traditional_rmse = results_df[results_df['Model'].isin(traditional_models)]['RMSE'].mean()
dl_rmse = results_df[results_df['Model'].isin(dl_models)]['RMSE'].mean()

if traditional_rmse < dl_rmse:
    better_approach = "traditional time series models"
    method_improvement = ((dl_rmse - traditional_rmse) / dl_rmse) * 100
else:
    better_approach = "deep learning approaches"
    method_improvement = ((traditional_rmse - dl_rmse) / traditional_rmse) * 100

print(f"4. Overall, {better_approach} performed better on this dataset (by {method_improvement}% improvement)")

# Calculate statistical significance if possible

```

```

print("\nConclusion:")
print("Based on our comprehensive analysis of four forecasting models (ARIMA, Proph

print(f"- Consumer spending is predictable using time series forecasting methods, w
print(f"- {better_approach.capitalize()} showed superior performance for this speci
print(f"- Training on {better_data} yields more accurate forecasts, suggesting that
print(f"- The 30-day forecasts from all models show similar trends, increasing confi

print("\nThis analysis successfully addresses Research Question 1: 'Can we predict
print("The answer is affirmative, with quantifiable accuracy metrics demonstrating

```

=====

FINAL SUMMARY AND INTERPRETATION

=====

Best Model for Full Dataset:

Model: LSTM

RMSE: 0.004752

MAE: 0.003622

Best Model for Recovery Phase Dataset:

Model: ARIMA

RMSE: 0.008813

MAE: 0.007297

Key Findings and Interpretation:

1. The full historical dataset generally produced better forecasts across models (by 9.50% in RMSE).
2. The LSTM model performed best on the full dataset with RMSE of 0.004752.
3. The ARIMA model performed best on the recovery phase data with RMSE of 0.008813.
4. Overall, traditional time series models performed better on this dataset (by 34.8 2% in RMSE).

Conclusion:

Based on our comprehensive analysis of four forecasting models (ARIMA, Prophet, LSTM, and Transformer) applied to consumer spending data, we can conclude that:

- Consumer spending is predictable using time series forecasting methods, with the best model (LSTM) achieving an RMSE of 0.004752.
- Traditional time series models showed superior performance for this specific economic indicator.
- Training on full historical dataset yields more accurate forecasts, suggesting that including pre-recovery patterns improves model performance.
- The 30-day forecasts from all models show similar trends, increasing confidence in the overall direction of future consumer spending patterns.

This analysis successfully addresses Research Question 1: 'Can we predict future consumer spending using time series forecasting models?'

The answer is affirmative, with quantifiable accuracy metrics demonstrating the effectiveness of these approaches.

In []:

In []:

In []:

In []:

In []:

In []: